

Django ORM crash test



Django ORM crash test

Andrii Soldatenko
9 April 2017
Italy, Otto

 @a_soldatenko

Andrii Soldatenko

- Senior Python Developer at



- CTO at



- Co-organizer PyCon Belarus 2017



- Speaker at many PyCons and open source contributor

- blogger at <https://asoldatenko.com>

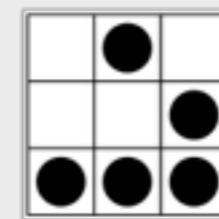


Table of contests:

- Object of testing
- SQL compare to Django ORM
- Crash tests and Django problems
- Conclusion



Andrew Svetlov
asvetlov

Python core developer

 @a_soldatenko



Paweł Piotr Przeradowski [Follow](#)

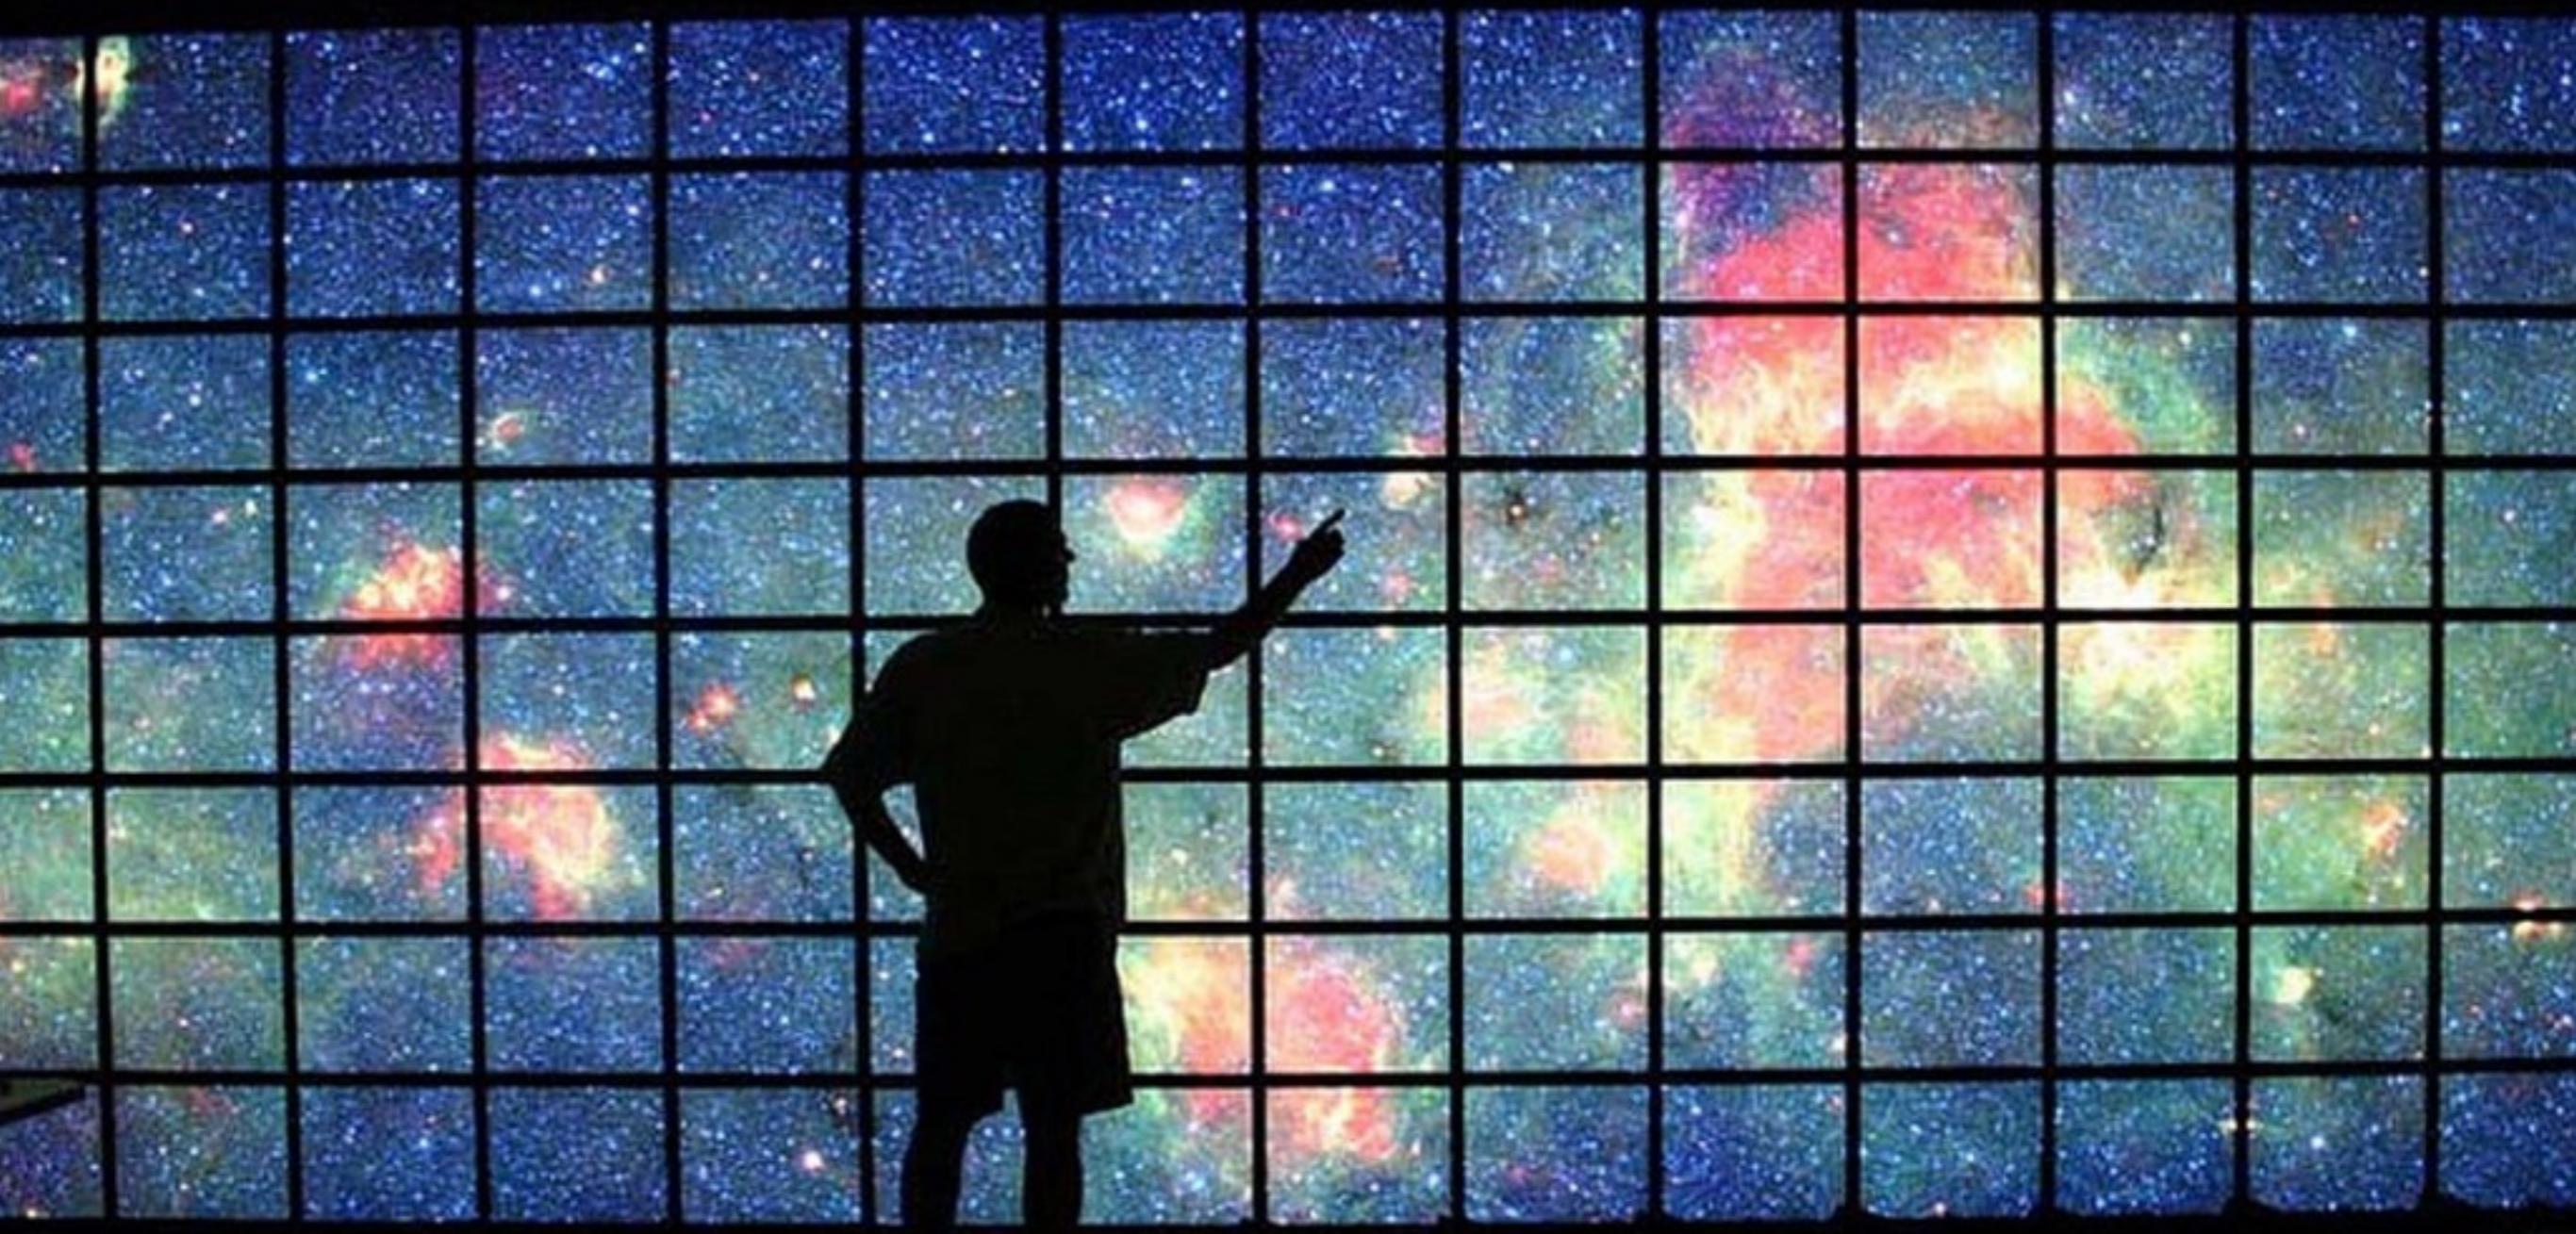
Python programmer, Web developer, Open source lover, Linux user, amateur photographer

4 days ago · 6 min read

A million requests per second with Python

<https://medium.freecodecamp.com/million-requests-per-second-with-python-95c137af319#.m9nhe1p2w>

Why?



cat VERSION.md

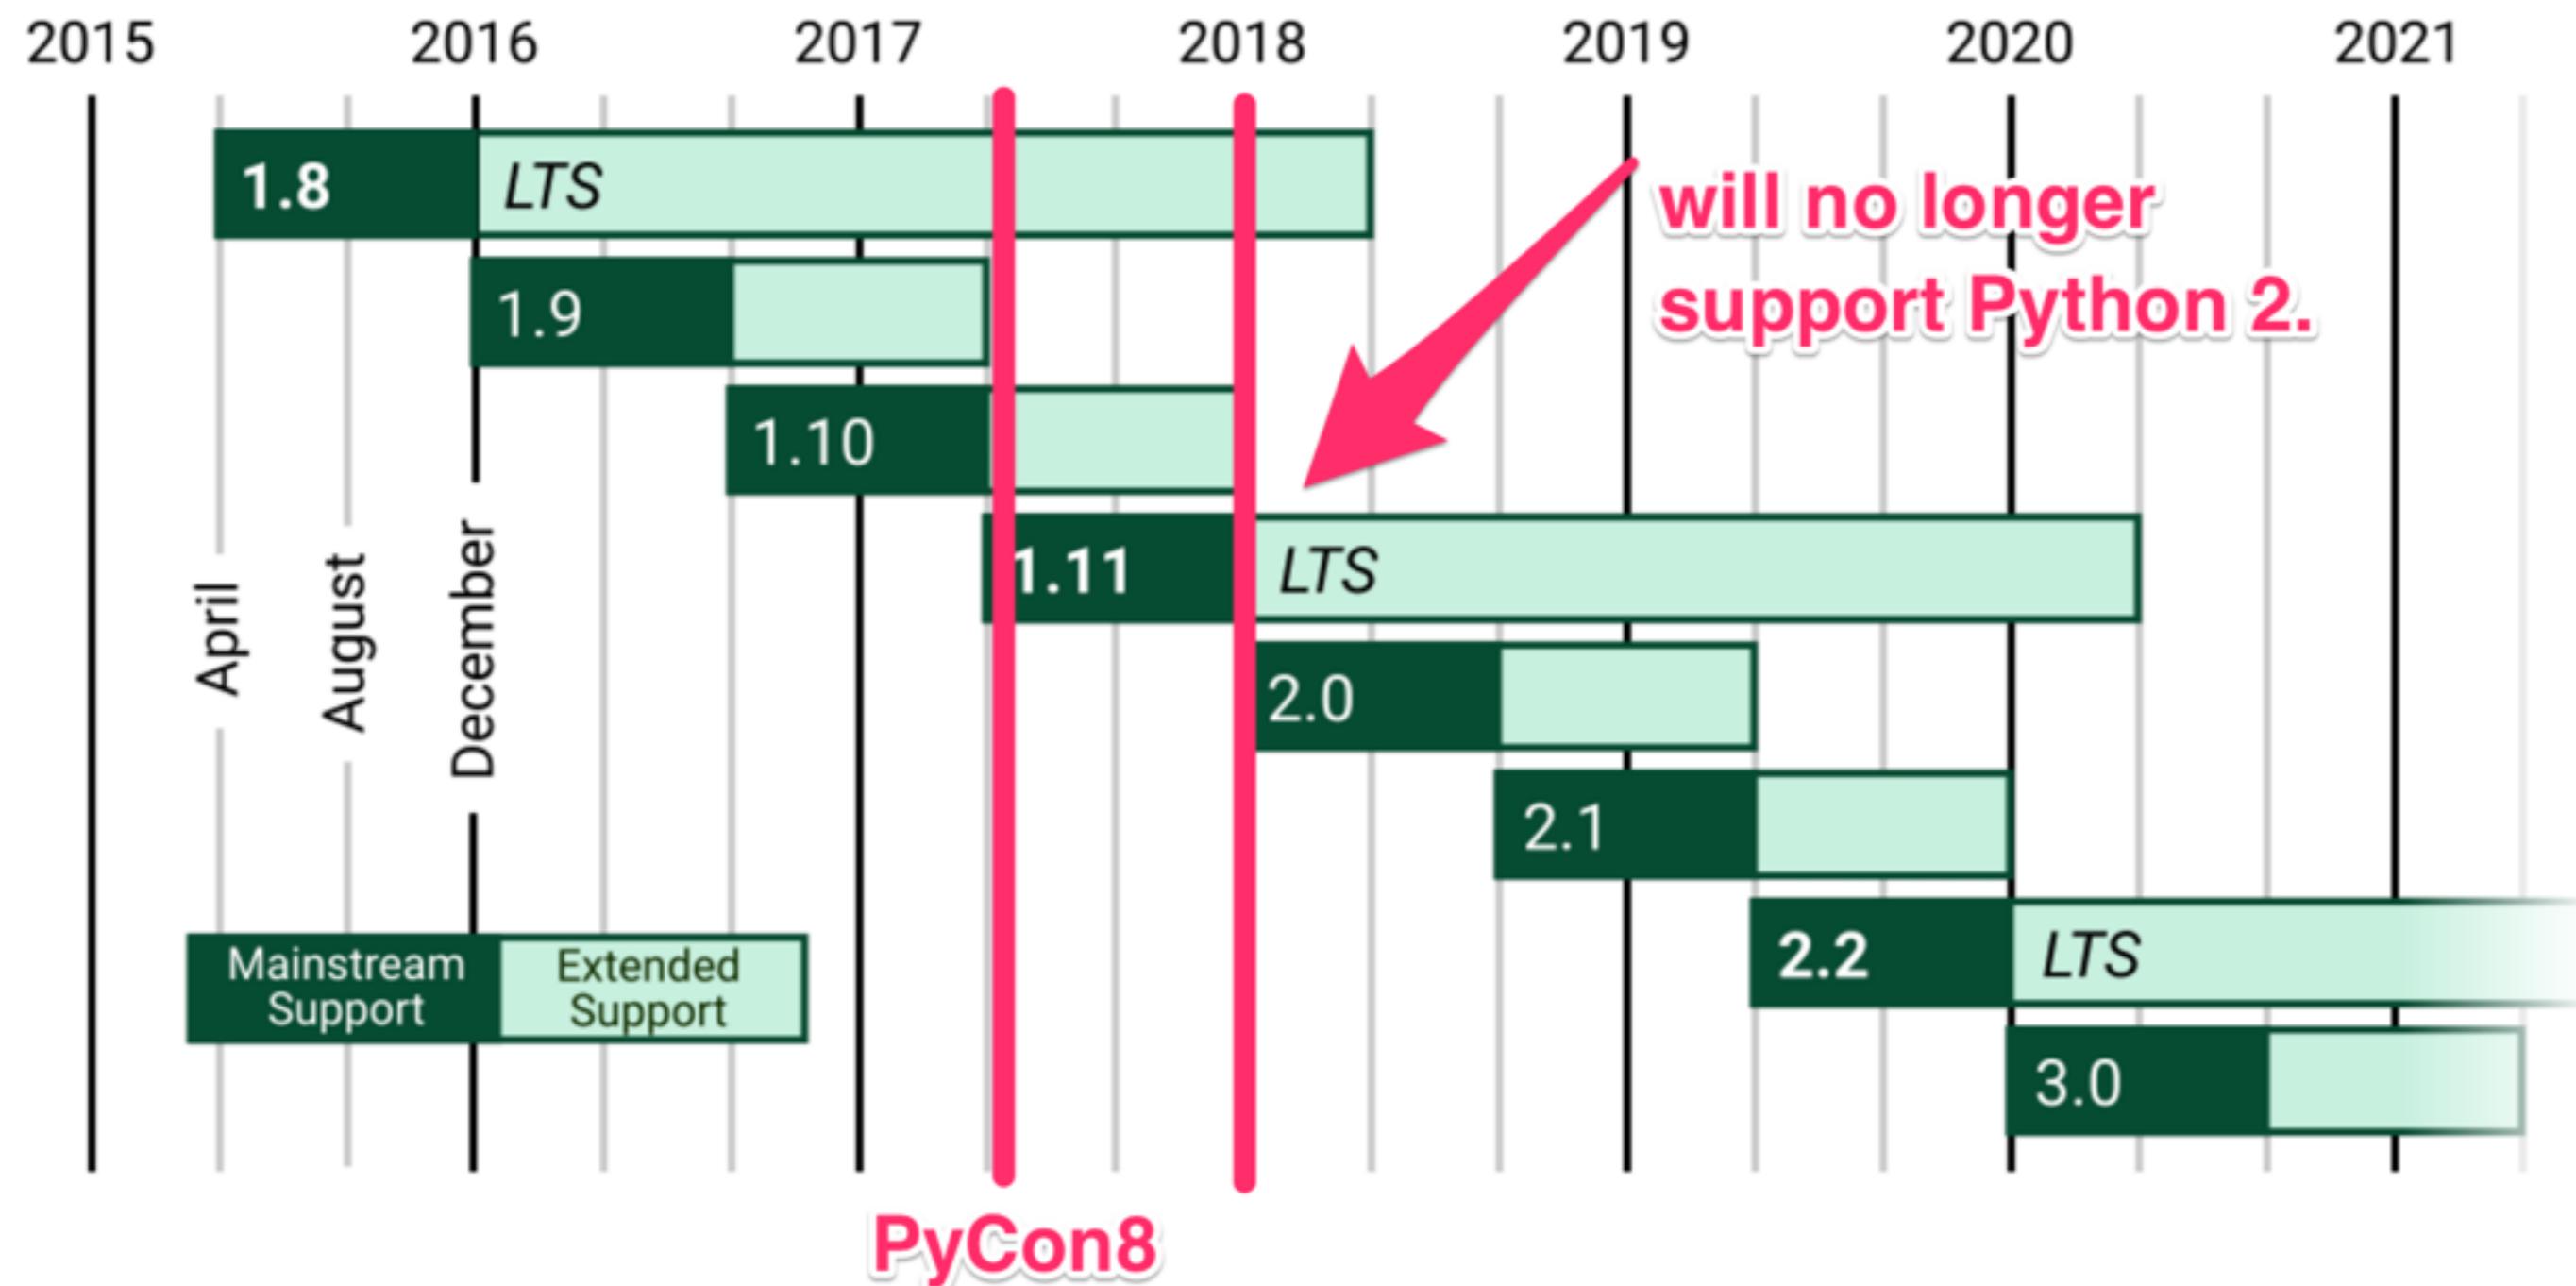
Python 3.6.1

Django 1.11

PostgreSQL 9.6.2

Docker 1.13

Django and Python 2



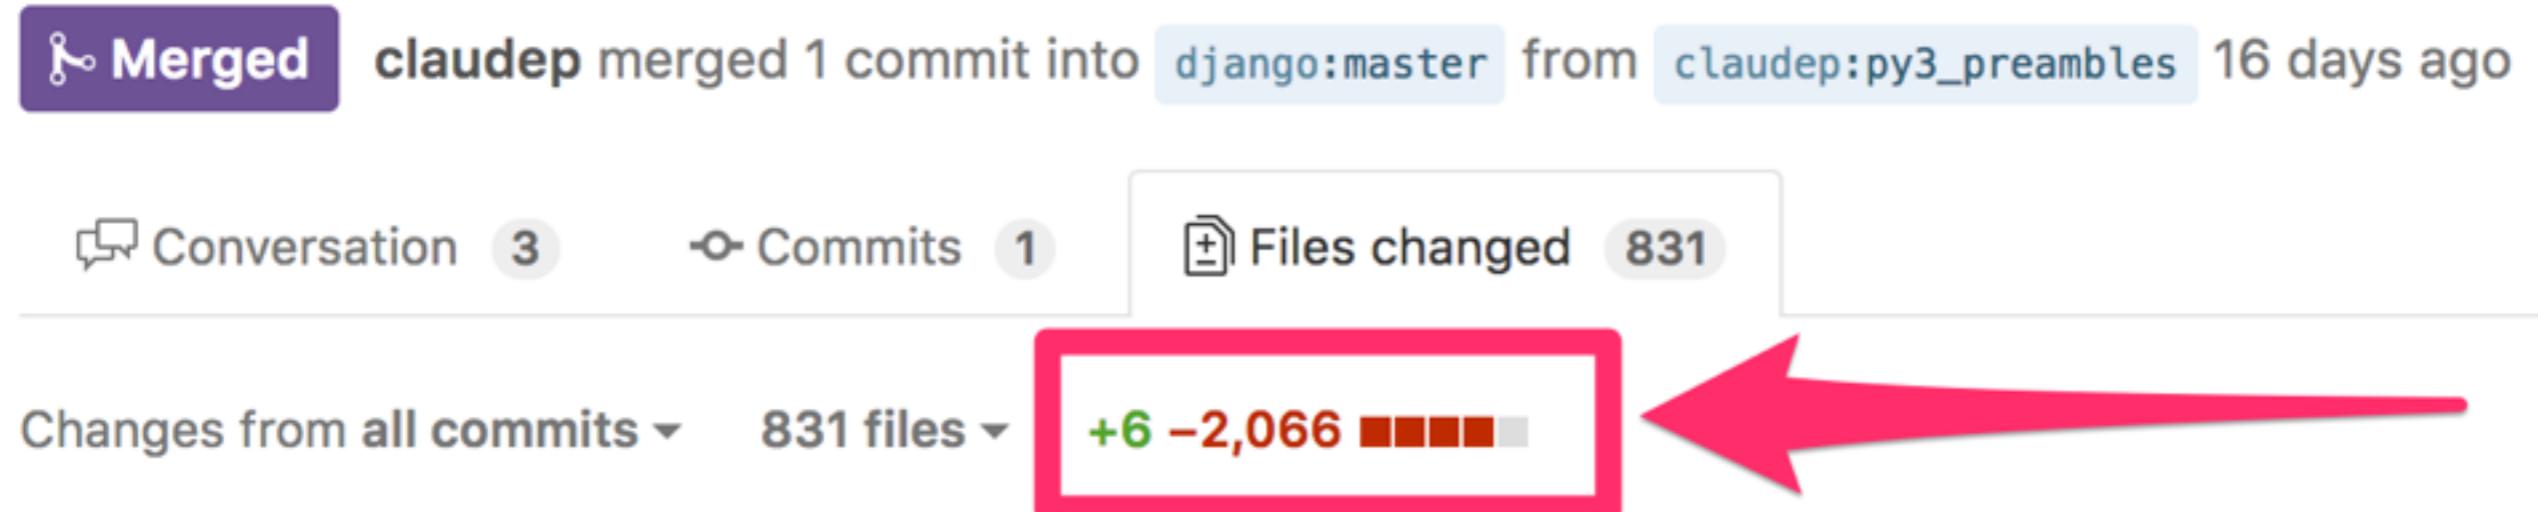
Remove Python 2 support

Refs #23919 -- Removed encoding preambles and
#7867

Merged claudep merged 1 commit into django:master from claudep:py3_preambles 16 days ago

Conversation 3 Commits 1 Files changed 831

Changes from all commits ▾ 831 files ▾ +6 -2,066



<https://github.com/django/django/pull/7867>

Real world example

ConsumerAffairs
Company



consumeraffairs.com

Similar web of CA

Engagement

Total Visits

7.10M

▼ 11.41%

⌚ Avg. Visit Duration

00:01:25

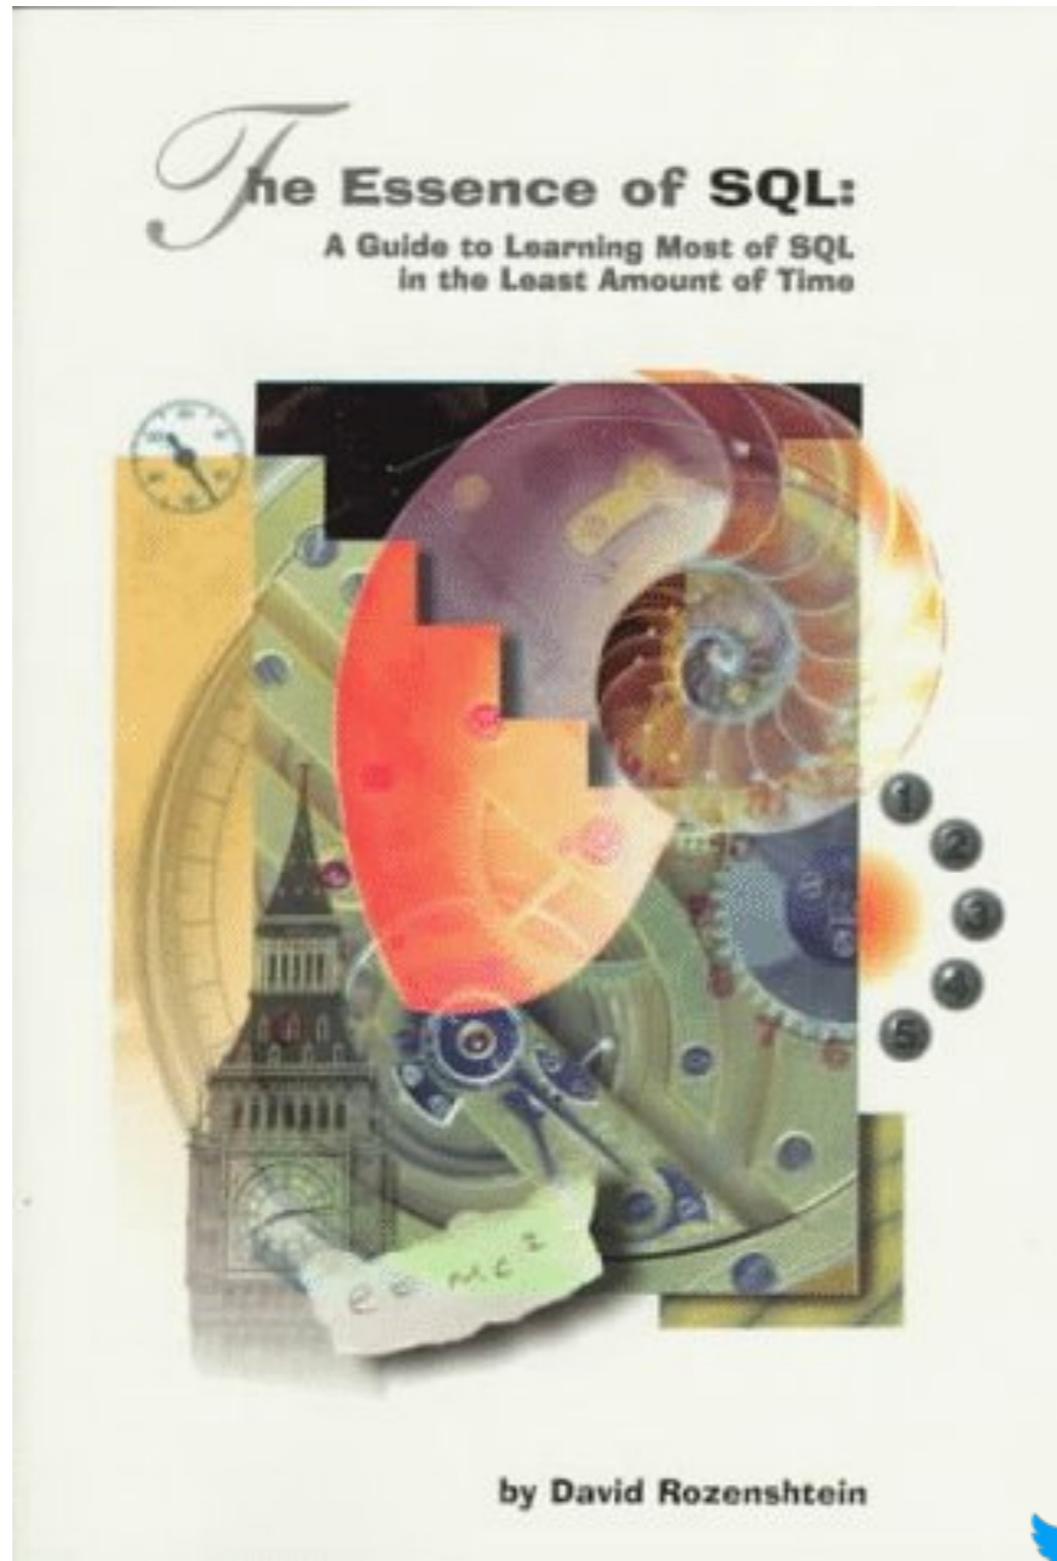
监听页面

1.75

↗ Bounce Rate

77.56%

Rozenshtein tests



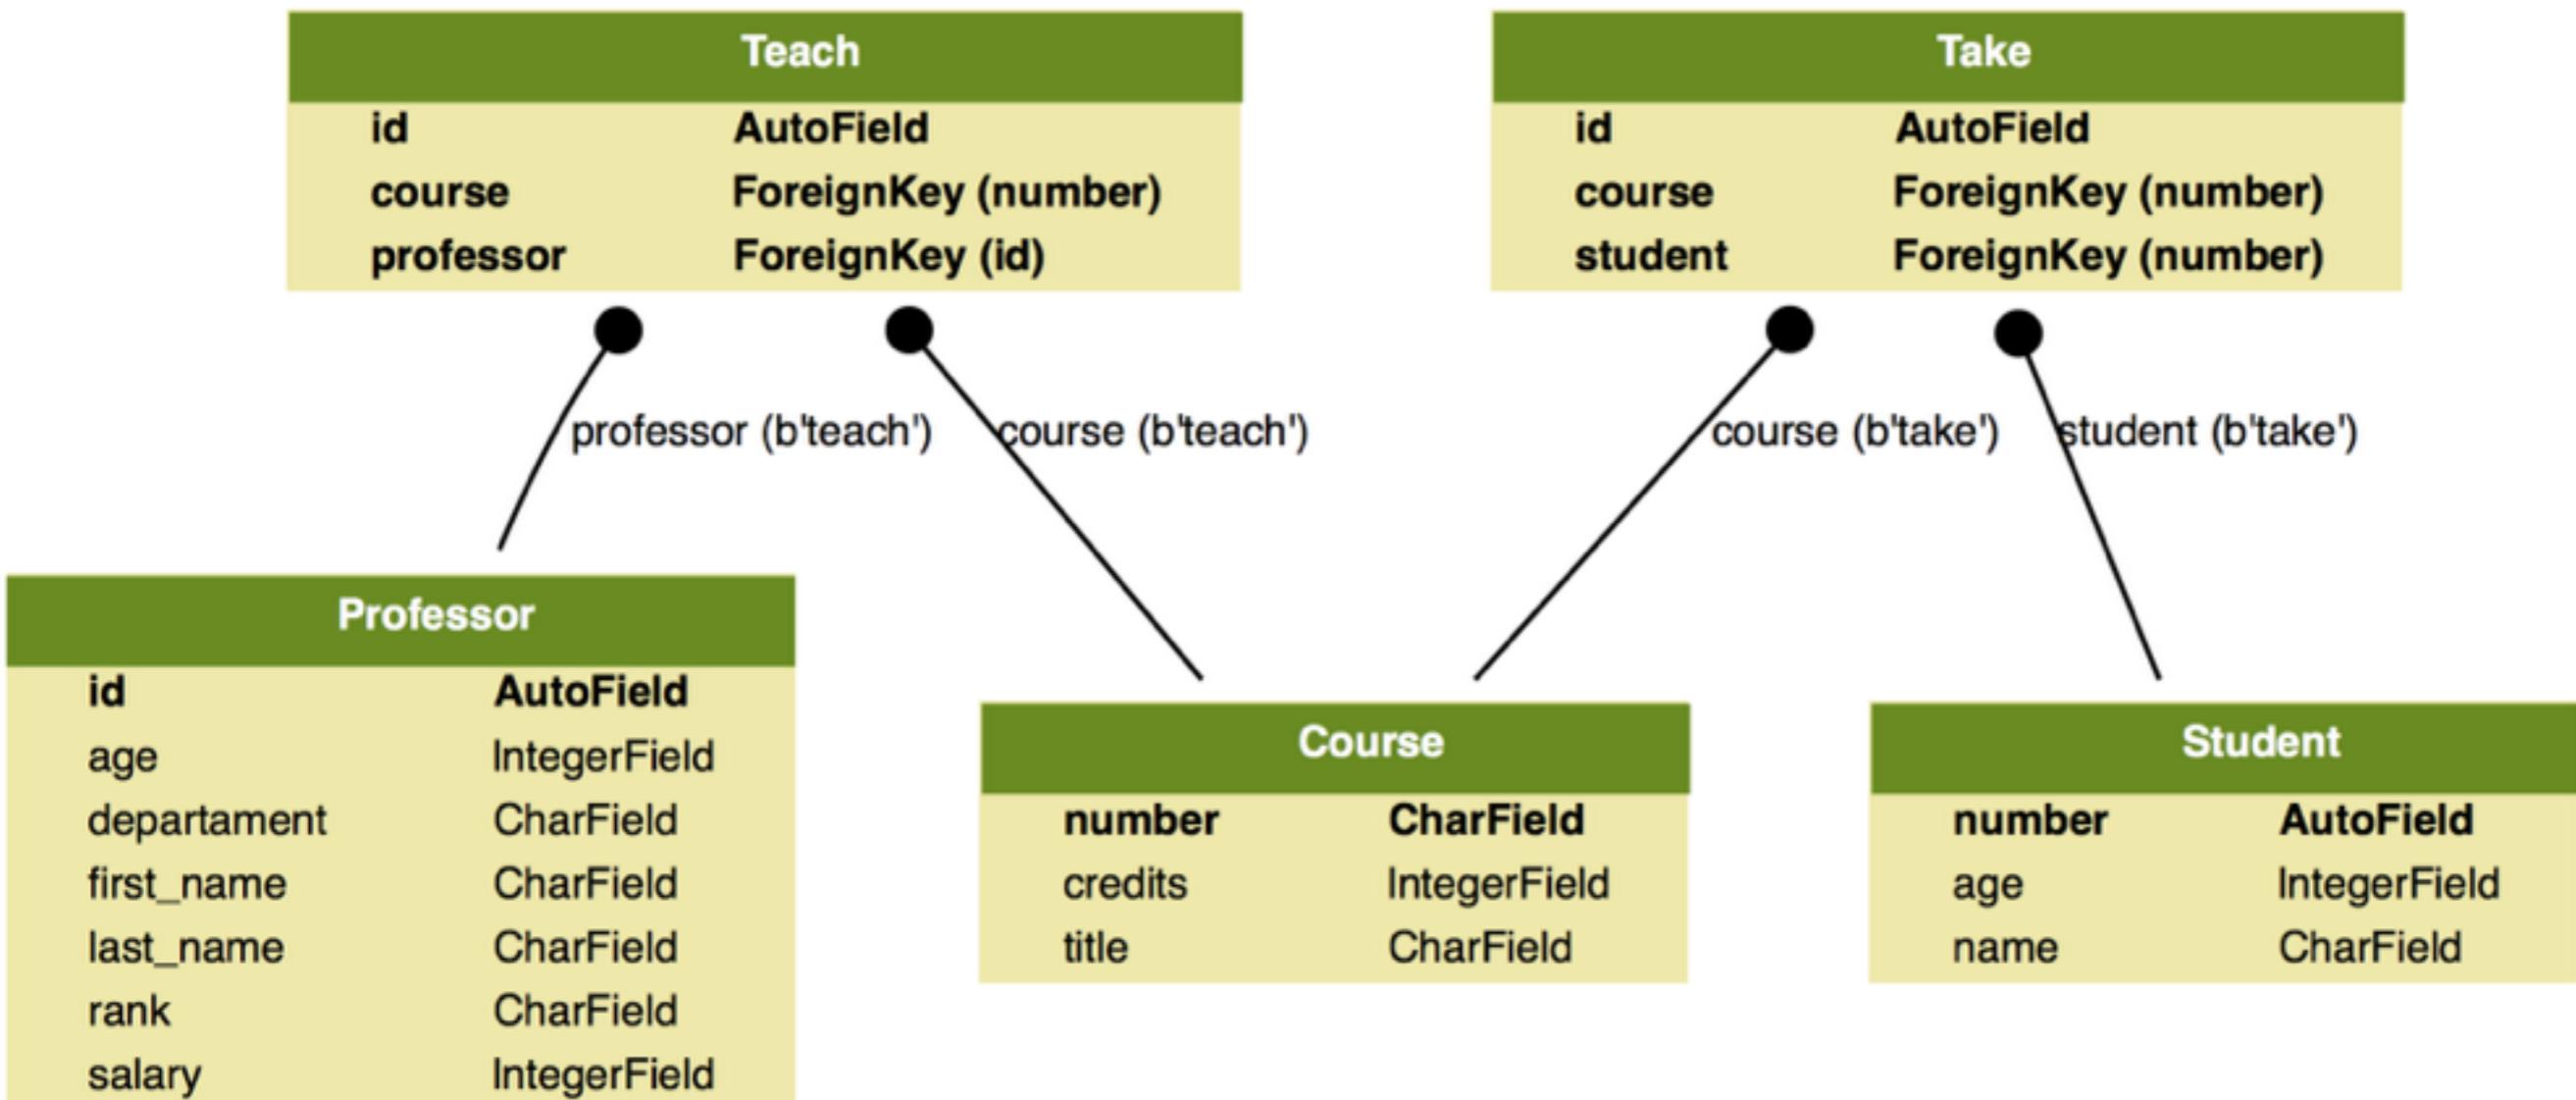
by David Rozenshtein

 @a_soldatenko

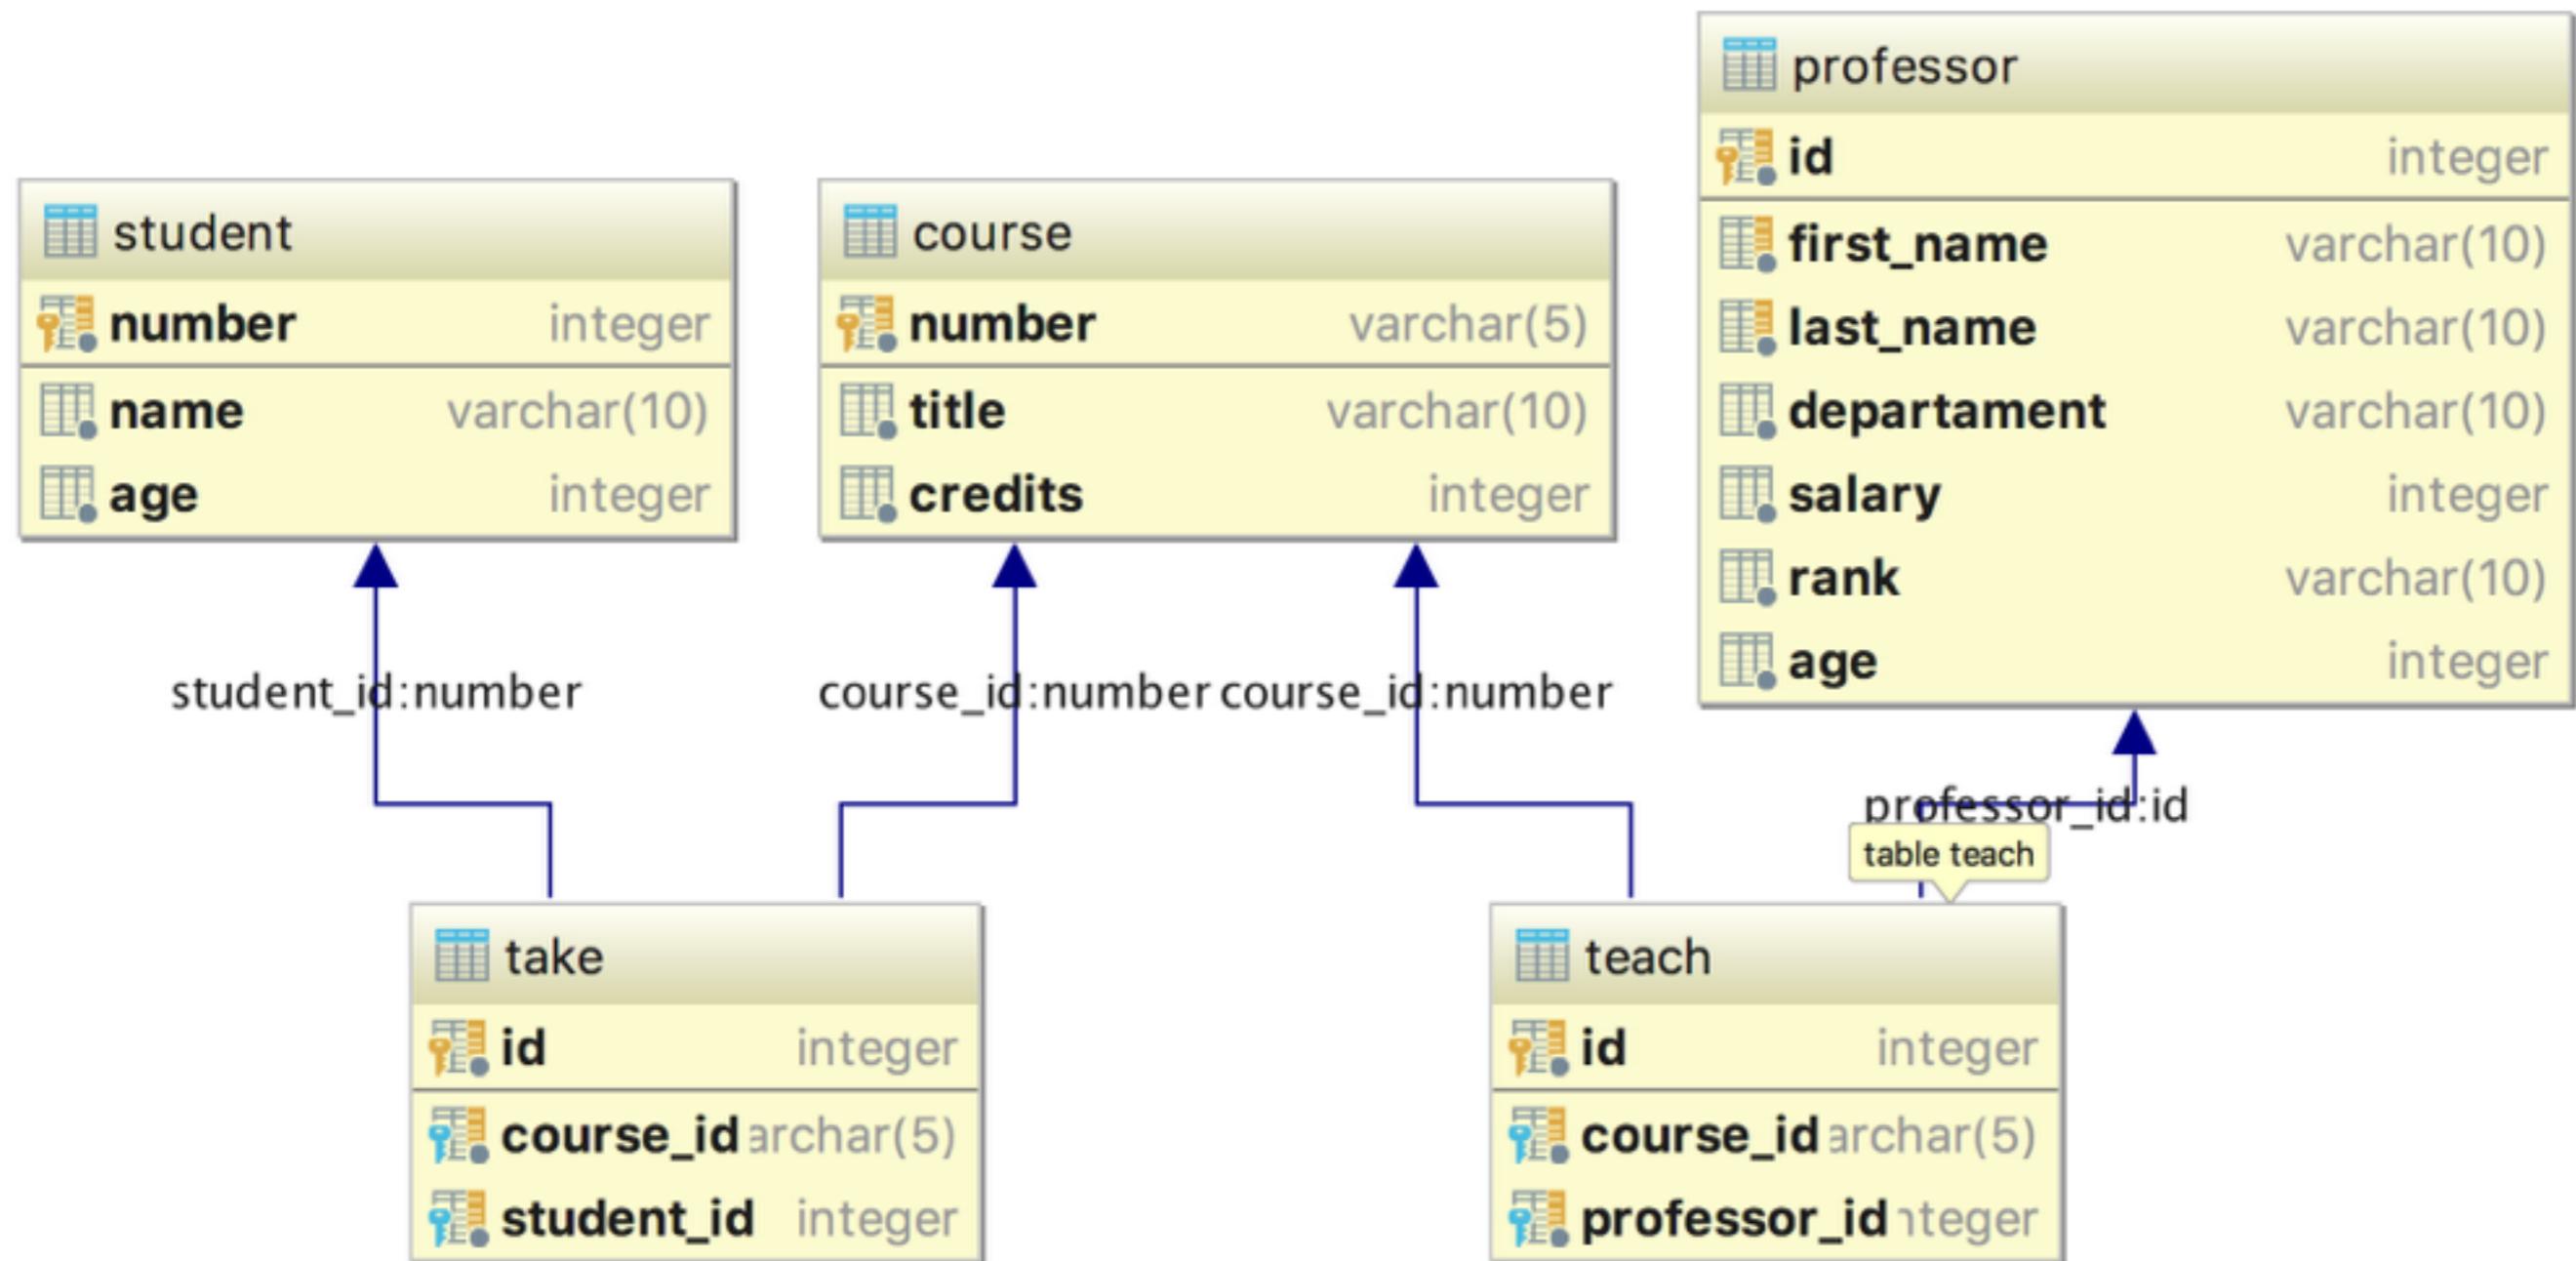
Django crash tests project

[https://github.com/andriisoldatenko/django-
orm-crash-test](https://github.com/andriisoldatenko/django-orm-crash-test)

University Database



University Database



A Bird's Eye View of SQL

"A Relational Model of Data for Large Shared Data Banks." - Edgar F. Codd's relational model 1970

DDL

```
CREATE TABLE course (
    number varchar(5) NOT NULL PRIMARY KEY,
    title varchar(10) NOT NULL,
    credits integer NOT NULL
);
```

```
ALTER TABLE teach ADD CONSTRAINT
"teach_course_id_bebbbd64_fk_course_number"
FOREIGN KEY (course_id) REFERENCES course
("number")
DEFERRABLE INITIALLY DEFERRED;
```

DML

SELECT ... **FROM** ... WHERE ...

INSERT INTO ... VALUES ...

UPDATE ... **SET** ... WHERE ...

DELETE **FROM** ... WHERE ...

PostgreSQL Hackers

<https://www.postgresql.org/list/pgsql-hackers/>

[https://wiki.postgresql.org/wiki/
So,_you_want_to_be_a_developer%3F](https://wiki.postgresql.org/wiki/So,_you_want_to_be_a_developer%3F)

Tom Lane



 @a_soldatenko

ORM under the hood

Primary keys problem

```
class Student(models.Model):  
    # id = models.AutoField(primary_key=True)  
    name = models.CharField(max_length=10)  
    age = models.IntegerField()
```



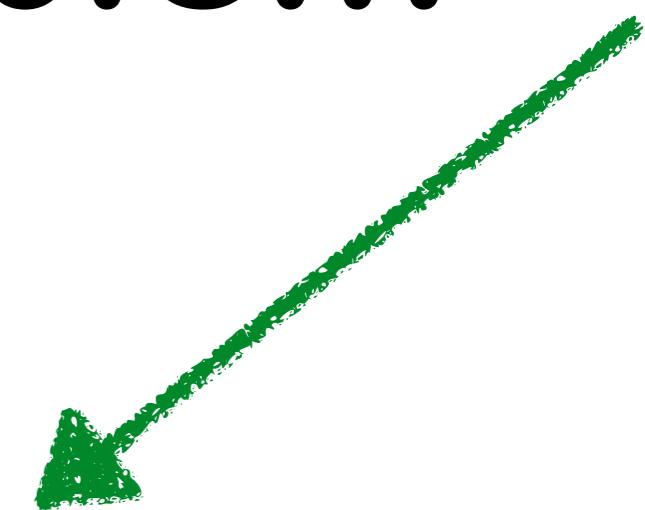
Primary keys problem

```
from django.db import models
```

```
class Student(models.Model):
    number = models.AutoField(primary_key=True)
    name = models.CharField(max_length=10)
    age = models.IntegerField()
```

```
def __str__(self):
    return self.name
```

```
class Meta:
    db_table = 'student'
```



Multi-Column Primary key problem

```
class Student(models.Model):
    id = models.AutoField(primary_key=True)
    number = models.AutoField(primary_key=True)
    name = models.CharField(max_length=10)
    age = models.IntegerField()

    new_class.add_to_class(obj_name, obj)
File "/Users/andrii/.pyenv/versions/django-orm-crash-test/lib/
python3.6/site-packages/django/db/models/base.py", line 328, in
add_to_class
    value.contribute_to_class(cls, name)
File "/Users/andrii/.pyenv/versions/django-orm-crash-test/lib/
python3.6/site-packages/django/db/models/fields/__init__.py", line 965,
in contribute_to_class
    assert not cls._meta.auto_field, "A model can't have more than one
AutoField."
AssertionError: A model can't have more than one AutoField.
```



@a_soldatenko

Multi-Column Primary keys problem PostgreSQL

```
CREATE TABLE student_w_two_primary (
    id INTEGER,
    number INTEGER,
    name VARCHAR(10) NOT NULL,
    age INTEGER NOT NULL,
    PRIMARY KEY (id, number)
);
```

Multi-Column Primary keys major issues

obj._meta.pk

backward for (content_type_id,
object_pk)

Admin URLs: "/app_label/
module_name/pk/"

Multi-Column Primary keys: Proposed Solutions

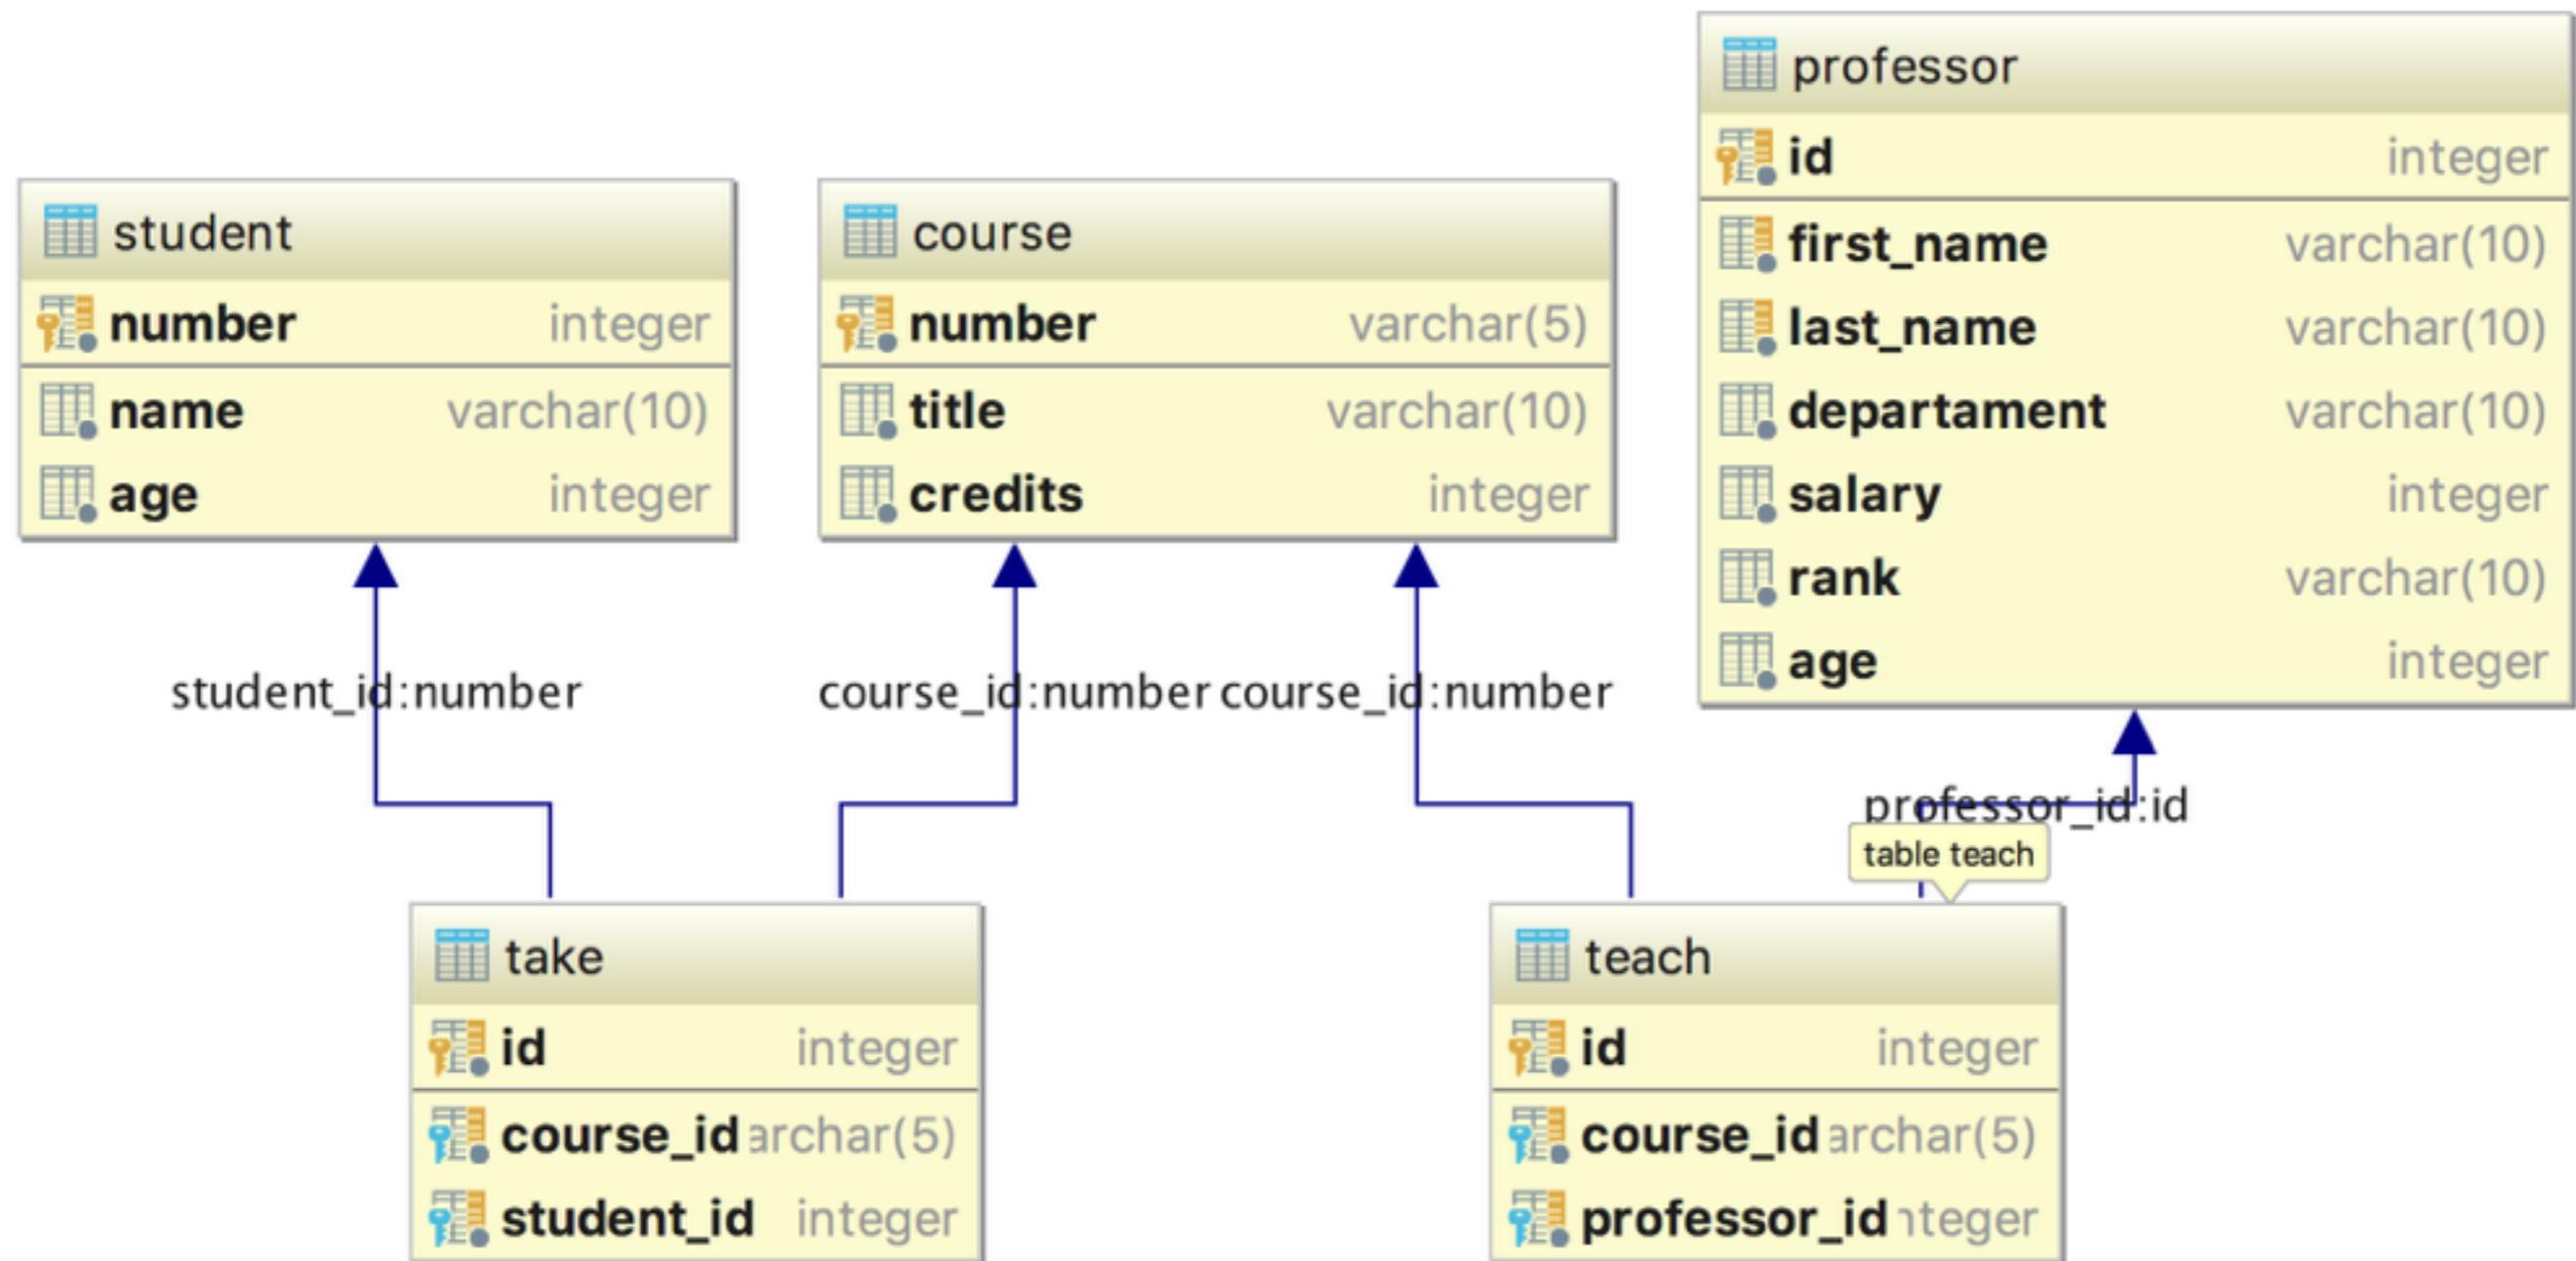
```
class Student(models.Model):
    id = models.BigIntegerField()
    number = models.AutoField(primary_key=True)
    name = models.CharField(max_length=10)
    age = models.IntegerField()

class Meta:
    db_table = 'student'
    unique_together = ('id', 'number')
```

Multi-Column Primary keys: Proposed Solutions

```
ALTER TABLE student ADD CONSTRAINT
    "student_id_number_345a5f6e_uniq"
UNIQUE ("id", "number");
```

University Database



How takes the course CS112?

```
Take.objects.filter(  
    course_number='CS112'  
).values('student')
```

```
SELECT "take"."student_id"  
FROM "take"  
WHERE "take"."course_id" = 'CS112';
```

What are student numbers and names of students who take CS112?

```
Take.objects.filter(  
    course_number='CS112'  
) .values('student_name', 'number')
```

```
SELECT "student"."name", "take"."student_id"  
FROM take  
INNER JOIN student  
ON ("take"."student_id" = "student"."number")  
WHERE "take"."course_id" = 'CS112';
```

Who takes the course CS112 or CS114?

```
Take.objects.filter(  
    course_number_in=['CS112', 'CS114'])  
.values('student_name', 'number')
```

```
SELECT "student"."name"  
FROM "take"  
INNER JOIN "student" ON  
    ("take"."student_id" = "student"."number")  
WHERE "take"."course_id" IN ('CS112', 'CS114')
```

Who takes the course CS112 or CS114?

```
from django.db.models import Q
```

```
Take.objects.filter(  
    Q(course_number='CS112') |  
    Q(course_number='CS114'))  
.values('student_name')
```

```
SELECT "student"."name"  
FROM "take"  
INNER JOIN "student" ON ("take"."student_id" =  
"student"."number")  
WHERE ("take"."course_id" = 'CS112'  
    OR "take"."course_id" = 'CS114');
```



@a_soldatenko

Who takes both courses CS112 and CS114?

```
Take.objects.filter(  
    course_number='CS112'  
).filter(course_number='CS114'  
).values('student_name')
```

```
SELECT "student"."name"  
FROM "take"  
INNER JOIN "student" ON ("take"."student_id" =  
"student"."number")  
WHERE ("take"."course_id" = 'CS112'  
    AND "take"."course_id" = 'CS114');
```



@a_soldatenko

Table inheritance problem



@a_soldatenko

Table inheritance problem

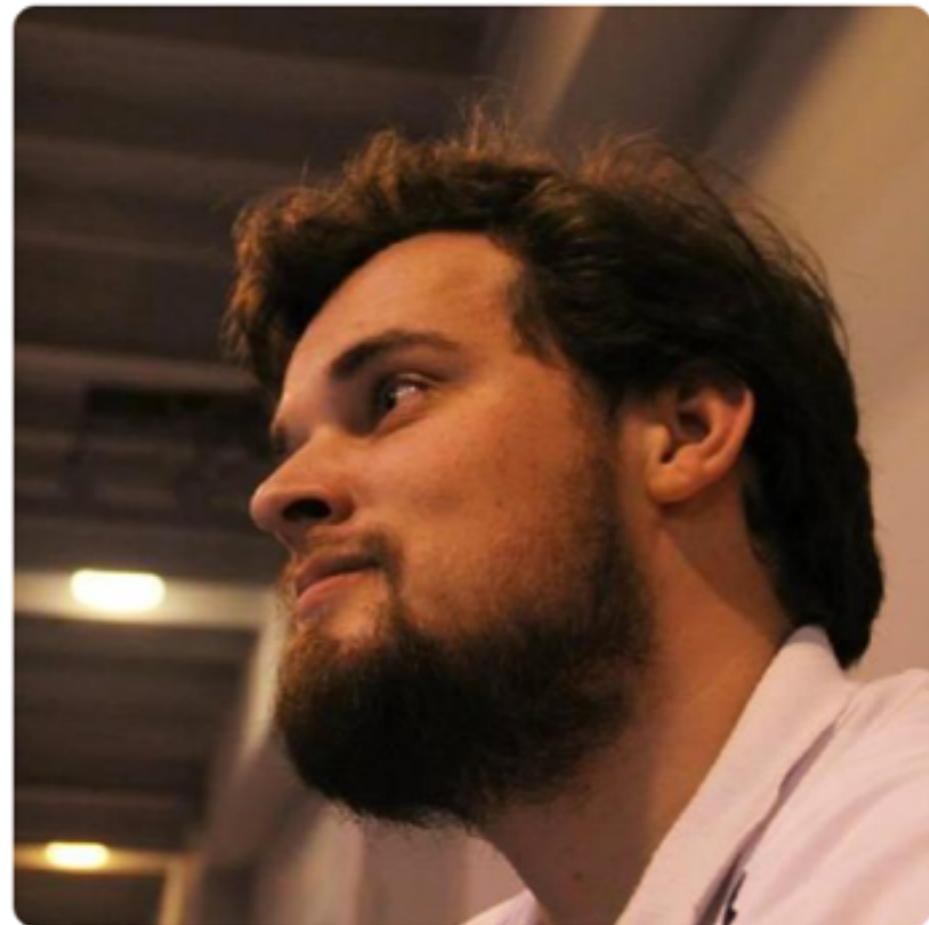
```
CREATE TABLE student
(
    number INTEGER PRIMARY KEY NOT NULL,
    name VARCHAR(10) NOT NULL,
    age INTEGER NOT NULL
);
```

```
CREATE TABLE professor
(
    departament VARCHAR(10) NOT NULL,
    salary INTEGER NOT NULL,
    rank VARCHAR(10) NOT NULL,
) INHERITS (student);
```

PostgreSQL table inheritance



...ultimately
rejected it as I'm
really not sure
how we can
handle it.



Marc Tamlyn
mjtamlyn

Enable SQL logging.v0

raw SQL queries Django is running?

```
# variant 0
```

```
In [1]: from django.db import connection
```

```
In [2]: Student.objects.all()
```

```
Out[2]: <QuerySet [<Student: AARON>, <Student: CHUCK>, <Student: DOUG>, ...]>
```

```
In [3]: connection.queries
```

```
Out[3]:
```

```
[{'sql': 'SELECT "student"."id",\n"student"."number", "student"."name",\n"student"."age" FROM "student" LIMIT 21',\n'time': '0.003'}]
```

Enable SQL logging.v1

raw SQL queries Django is running?

```
# variant 1
# pip install django-extensions
./manage.py shell_plus --print-sql
```

```
In [1]: student = Student.objects.first()
SELECT "student"."id", "student"."number",
"student"."name", "student"."age" FROM "student"
ORDER BY "student"."number" ASC LIMIT 1
```

```
Execution time: 0.002584s [Database: default]
```

Enable SQL logging.v2

raw SQL queries Django is running?

```
# variant 3
```

```
# pip install django-debug-toolbar
```

```
./manage.py debugsqlshell
```

```
In [2]: from university.models import Student
```

```
In [3]: Student.objects.all()
```

```
Out[3]: SELECT "student"."id",
          "student"."number",
          "student"."name",
          "student"."age"
```

```
FROM "student" LIMIT 21 [3.10ms]
```

```
<QuerySet [<Student: AARON>, <Student: CHUCK>, ]>
```

Enable SQL logging.v3

```
# variant 3
```

```
In [2]: from university.models import Student
```

```
In [4]: import logging
```

```
...: l =
```

```
logging.getLogger('django.db.backends')
```

```
...: l.setLevel(logging.DEBUG)
```

```
...: l.addHandler(logging.StreamHandler())
```

```
In [6]: Student.objects.first()
```

```
(0.003) SELECT "student"."id",
```

```
"student"."number", "student"."name",
```

```
"student"."age" FROM "student" ORDER BY
```

```
"student"."number" ASC LIMIT 1; args=()
```

```
Out[6]: <Student: AARON>
```

 @a_soldatenko

Enable SQL logging.v4

```
In [1]: from django.test.utils import CaptureQueriesContext
In [6]: with CaptureQueriesContext(connection) as context:
...
print(Take.objects.filter(Q(course_number='CS112')).values(
    'student_name'))
...
<QuerySet [{ 'student_name': 'AARON'}, {'student_name': 'CHUCK'}, {'student_name': 'DOUG'}, {'student_name': 'MAGGIE'}]>
In [8]: print(context.captured_queries)
[{'sql': 'SELECT "student"."name" FROM "take" INNER JOIN "student" ON ("take"."student_id" = "student"."number") WHERE "take"."course_id" = \'CS112\' LIMIT 21', 'time': '0.005'}]
```

Enable SQL logging.v5

Transaction trace Track as Key Transaction 

[Delete this trace](#)

Dec 9, '16 8:38 pm 4,890 ms 301 ms (6.17%) 77.3 ms (1.58%)

TRACE TIME RESP. TIME USER CPU BURN SYSTEM CPU BURN

Summary Trace details **Database queries**

Show fast queries (< 5ms)

Total duration	Call count	Query
10 ms	7	SELECT fc.*FROM field_config fc WHERE (fc.field_name = :db_condition_placeholder_?) AND (fc.active = :db_condition_placeholder_?) AND (fc.storage_active = :db_condition_placeholder_?) AND (fc.deleted = :db_condition_placeholder_?)
98 ms	3	DELETE FROM variable WHERE (name = :db_condition_placeholder_?)
4,340 ms	1	INSERT INTO semaphore (name, value, expire) VALUES (:db_insert_placeholder_?, :db_insert_placeholder_?, :db_insert_placeholder_?)

Query details 

Duration	4,344 ms
----------	----------

Stack trace

```
in PDOStatement::execute called at /mnt/www/html/newrelic/docroot/includes/database/database.inc (2171)
in DatabaseStatementBase::execute called at /mnt/www/html/newrelic/docroot/includes/database/database.inc (683)
in DatabaseStatementBase::execute called at /mnt/www/html/newrelic/docroot/includes/database/database.inc (683)
```



@a_soldatenko

SQL string concatenation

```
Student.objects.all().update(age=F('age')+10)
```

```
UPDATE "student"  
SET "age" = ("student"."age" + 10)
```

SQL AS column aliases

```
In [33]: Course.objects.extra(select={'id':  
'number'}) .values()
```

```
Out[33]: SELECT (number) AS "id",  
    "course"."number",  
    "course"."title",  
    "course"."credits"  
  
FROM "course" LIMIT 21
```

```
<QuerySet [{  
    'id': 'CS112', 'number': 'CS112',  
    'title': 'PHYSICS', 'credits': 4},  
    {'id': 'CS113',  
    'number': 'CS113', 'title': 'CALCULUS', 'credits':  
    4},  
    {'id': 'CS114', 'number': 'CS114', 'title':  
    'HISTORY', 'credits': 4}]>
```

SQL Functions

Concat

Greatest

and more

Migration problem

```
x ./manage.py sqlmigrate university 0002
BEGIN;

-- 
-- Add field id to student
--

ALTER TABLE "student" ADD COLUMN "id" bigint DEFAULT 1 NOT
NULL;
ALTER TABLE "student" ALTER COLUMN "id" DROP DEFAULT;
--

-- Alter unique_together for student (1 constraint(s))
--

ALTER TABLE "student" ADD CONSTRAINT
"student_id_number_345a5f6e_uniq" UNIQUE ("id", "number");
COMMIT;
```

Full text search problem

EuroPython 2016

July 21, 2016

What is the best full text search engine for python?



EUROPYTHON
2016 Bilbao, 17-24 July

Full text search
problem

New in Django 1.10.

But since

PostgreSQL 8.3

django.contrib.postgres.search

In [15]:

```
Course.objects.annotate(search=SearchVector('title')) .filter(search='PHYSIC')
```

Out[15]: SELECT "course"."number",

"course"."title",

"course"."credits",

to_tsvector(COALESCE("course"."title", ''))

AS "search"

FROM "course"

WHERE to_tsvector(COALESCE("course"."title", ''))

@@ (plainto_tsquery('PHYSIC')) = true LIMIT 21

Execution time: 0.002756s [Database: default]

<QuerySet [<Course: CS112: PHYSICS>]>

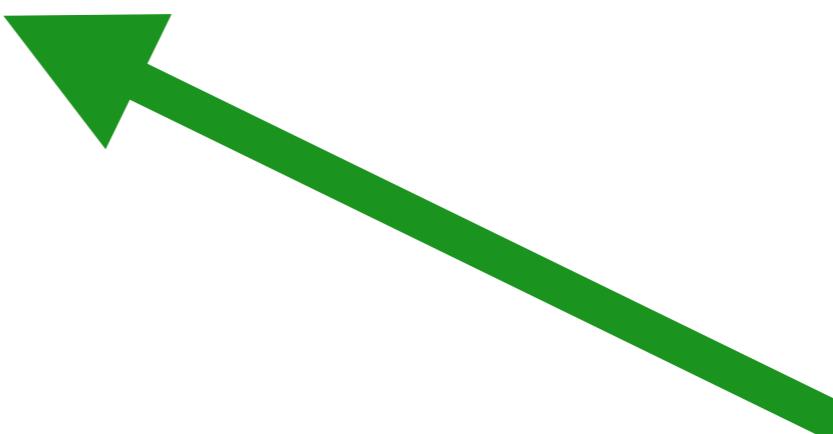
Two nuances with FTS in Django



@a_soldatenko

You should add `search_field` create FTS field

```
class Course(models.Model):  
    number = models.CharField(...)  
    title = models.CharField(..)  
    credits = models.IntegerField()  
    search_vector = SearchVectorField()
```



You should manually create trigger to update search_vector

```
CREATE TRIGGER tsvectorupdate BEFORE  
INSERT OR UPDATE  
ON course FOR EACH ROW EXECUTE  
PROCEDURE  
tsvector_update_trigger(search_vector,  
'pg_catalog.english', title);"
```

If you want to use different PG search functions

```
class TsQuerySearchQuery(SearchQuery):
    def as_sql(self, compiler, connection):
        params = [self.value]
        if self.config:
            config_sql, config_params =
compiler.compile(self.config)
            template = 'to_tsquery({}::regconfig,
%s)'.format(config_sql)
            params = config_params + [self.value]
        else:
            template = 'to_tsquery(%s)'
        if self.invert:
            template = '!!({})'.format(template)
    return template, params
```

And search query will looks like

```
Course.objects.filter(  
    search=TsQuerySearchQuery( 'PHYSIC:A' )  
)
```

cd django
git checkout 1.11.x

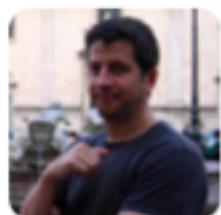
Subquery

Exists

OuterRef

<https://docs.djangoproject.com/en/1.11/ref/models/expressions/#django.db.models.Subquery>

Async database problem



mike bayer
@zzzeek

Following

the reason I ask is bc, I've tried for 2 days to get
asyncio *or* gevent to outperform threads on a
real RDBMS speed test, and I can't.

RETWEETS

4

LIKES

4



1:44 AM - 13 Feb 2015

5

4

4

<http://techspot.zenzeek.org/2015/02/15/asynchronous-python-and-databases/>

Performance problem: pgbench

```
x docker exec -it postgresql sudo -u postgres pgbench -i -s 15  
university_db  
creating tables...  
100000 of 1500000 tuples (6%) done (elapsed 0.77 s, remaining  
10.79 s)  
200000 of 1500000 tuples (13%) done (elapsed 2.37 s, remaining  
15.39 s)  
300000 of 1500000 tuples (20%) done (elapsed 4.13 s, remaining  
16.51 s)
```

set primary keys...
done.

<https://www.postgresql.org/docs/9.6/static/pgbench.html>

Final Thoughts

- Thanks David Rozenshtein for his tests
- django.contrib.postgres
- New crash-test SQLAlchemy
- github.com/andriisoldatenko/django-orm-crash-test

Thank You

andrii.soldatenko@toptal.com

https://asoldatenko.com

Questions



We are hiring



<https://www.toptal.com/#connect-fantastic-computer-engineers>