



EUROPYTHON
2016 Bilbao, 17-24 July

What is the best full text search engine for Python?

Andrii Soldatenko
[@a_soldatenko](https://twitter.com/a_soldatenko)

Agenda:

- Who am I?
- What is full text search?
- PostgreSQL FTS / Elastic / Whoosh / Sphinx
- Search accuracy
- Search speed
- What's next?

Andrii Soldatenko

- Backend Python Developer at
- CTO in Persollo.com
- Speaker at many PyCons and Python meetups
- blogger at <https://asoldatenko.com>



Preface



Linktom®

Text Search

→ cpython time grep -r -i 'OrderedDict' .

grep -r -i 'OrderedDict' **2.35s user 0.10s system 97% cpu 2.510 total**

→ cpython time ack OrderedDict

ack OrderedDict **1.74s user 0.14s system 96% cpu 1.946 total**

→ cpython time pss OrderedDict

pss OrderedDict **0.85s user 0.09s system 96% cpu 0.983 total**

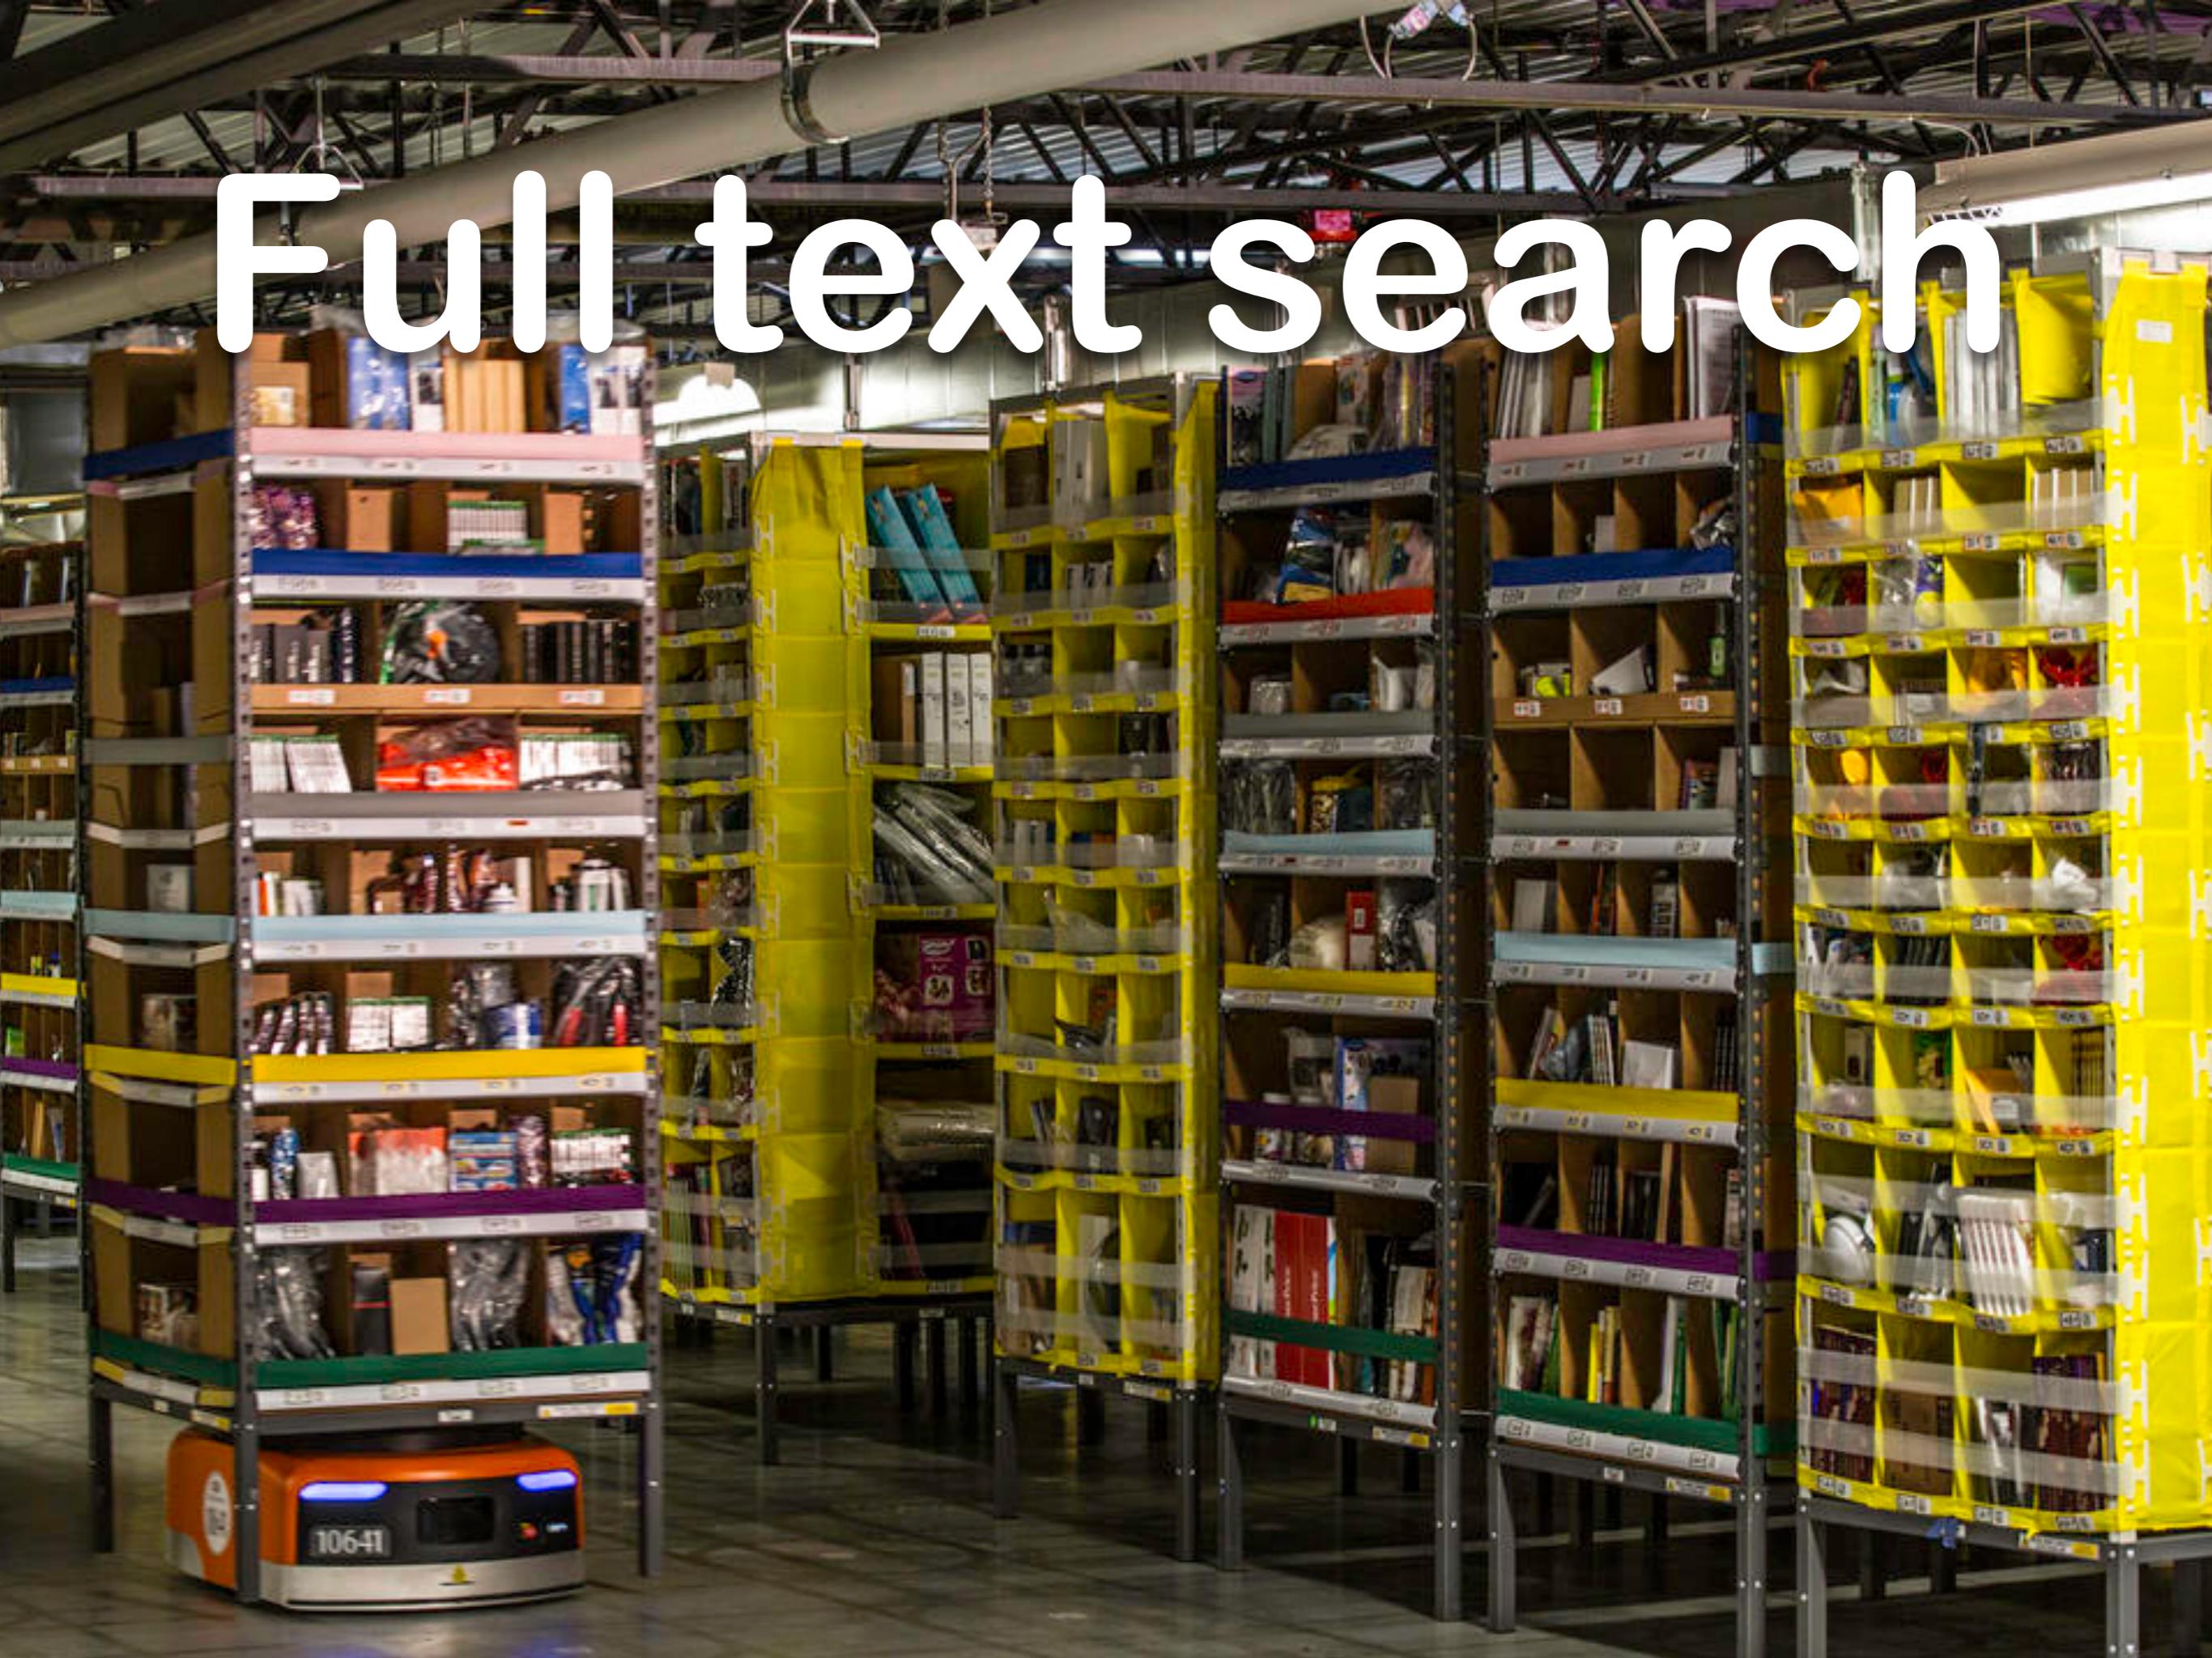
→ cpython time pt OrderedDict

pt OrderedDict **0.14s user 0.10s system 462% cpu 0.051 total**



Processor 2.5 GHz Intel Core i7
Memory 16 GB 1600 MHz DDR3

Full text search



Search index

Symbols

32gb Heap boundary, 642

A

ACID transactions, 545, 556

action, in bulk requests, 57, 69

ad hoc searches, 15

aggregations, 20, 417

 aggs parameter, 424

 and analysis, 483

 approximate, 457

 cardinality, 458

 percentiles, 462

 basic example

 adding a metric, 426

 adding extra metrics, 429

 buckets nested in other buckets, 427

buckets, 419

 combining buckets and metrics, 420

 metrics, 420

limiting memory usage, 487

 fielddata circuit breaker, 490

 fielddata size, 488

 moitoring fielddata, 489

managing efficient memory usage, 507

nested, 567

 reverse_nested aggregation, 568

operating alongside search requests, 418

preventing combinatorial explosions, 500

 depth-first versus breadth-first, 502

returning empty buckets, 439

scoping, 445

 global bucket, 447

Significant Terms, 471

significant terms

Simple sentences

1. The quick brown fox jumped over the lazy dog
2. Quick brown foxes leap over lazy dogs in summer

Inverted index

1	Term	Doc_1	Doc_2
2		-----	
3	Quick		X
4	The	X	
5	brown	X	X
6	dog	X	
7	dogs		X
8	fox	X	
9	foxes		X
10	in		X
11	jumped	X	
12	lazy	X	X
13	leap		X
14	over	X	X
15	quick	X	
16	summer		X
17	the	X	
18		-----	

Inverted index

1	Term	Doc_1	Doc_2
2	-----		
3	brown		X X
4	quick		X
5	-----		
6	Total	2	1

Inverted index: normalization

Term	Doc_1	Doc_2
Quick	X	X
The	X	
brown	X	X
dog	X	
dogs	X	
fox	X	
foxes	X	
in	X	
jumped	X	
lazy	X X	
leap	X	
over	X X	
quick	X	
summer	X	
the	X	

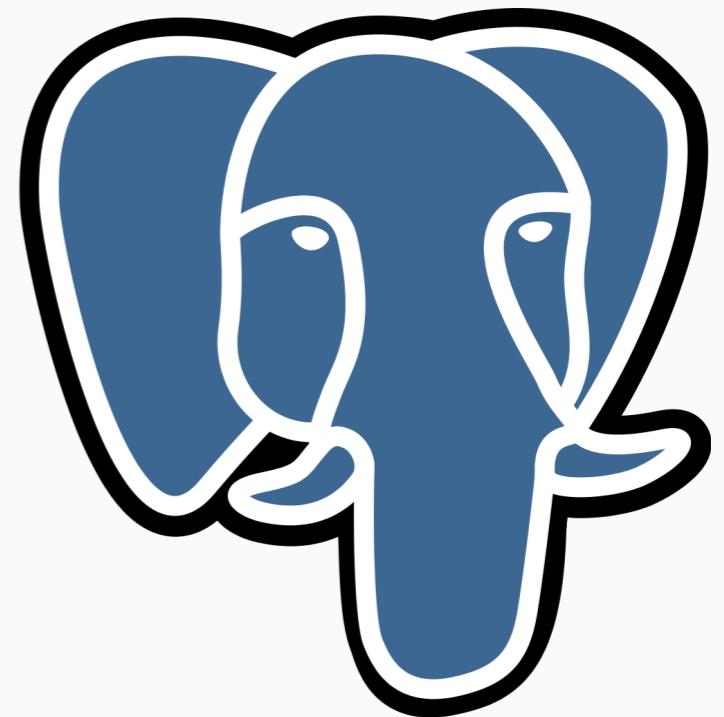
→

Term	Doc_1	Doc_2
brown	X X	
dog	X X	
fox	X X	
in	X	
jump	X X	
lazy	X X	
over	X X	
quick	X X	
summer	X	
the	X X	

Search Engines

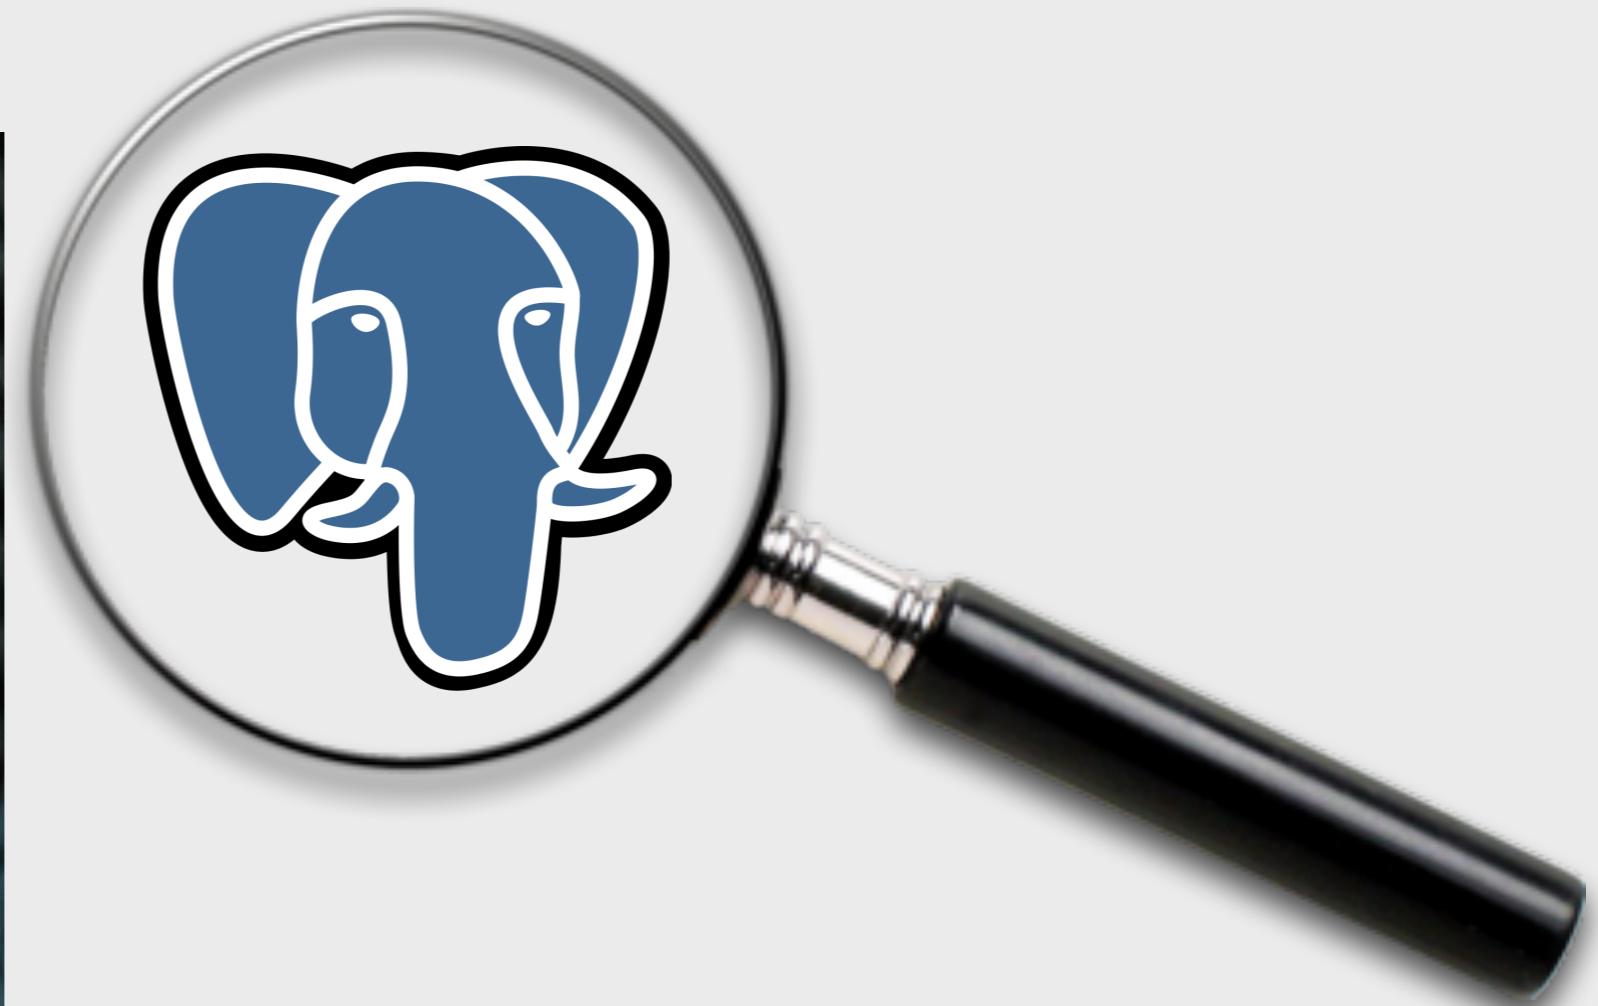


elastic



PostgreSQL

Full Text Search



support from version 8.3

PostgreSQL

Full Text Search

```
SELECT to_tsvector('text') @@  
to_tsquery('query');
```

Simple is better than complex. - by import this

Do PostgreSQL FTS without index

```
SELECT 'python bilbao 2016'::tsvector @@ 'python &  
bilbao'::tsquery;
```

?column?

t

(1 **row**)

Do PostgreSQL FTS with index

**CREATE INDEX name ON table USING GIN
(column);**

**CREATE INDEX name ON table USING
GIST (column);**

PostgreSQL FTS: Ranking Search Results

ts_rank() → float4 – based on the frequency of their matching lexemes

ts_rank_cd() → float4 – cover density ranking for the given document vector and query

PostgreSQL FTS Highlighting Results

```
SELECT ts_headline('english',
    'python conference 2016',
    to_tsquery('python & 2016'));
```

ts_headline

```
<b>python</b> conference <b>2016</b>
```

Stop Words

until not in
few or him ours it off
does their your my he be
as do we i ourselves a
she myself am
through what himself them
and

PostgreSQL FTS Stop Words

```
SELECT to_tsvector('in the list of stop words');
```

```
to_tsvector
```

```
'list':3 'stop':5 'word':6
```

PG FTS and Python

- Django 1.10 django.contrib.postgres.search
- djorm-ext-pgfulltext
- sqlalchemy

PostgreSQL FTS integration with django orm

```
from djorm_pgfulltext.models import SearchManager
from djorm_pgfulltext.fields import VectorField
from django.db import models

class Page(models.Model):
    name = models.CharField(max_length=200)
    description = models.TextField()

    search_index = VectorField()

    objects = SearchManager(
        fields = ('name', 'description'),
        config = 'pg_catalog.english', # this is default
        search_field = 'search_index', # this is default
        auto_update_search_field = True
    )
```

<https://github.com/linuxlewis/djorm-ext-pgfulltext>

For search just use search method of the manager

```
>>> Page.objects.search("documentation & about")
[<Page: Page: Home page>]

>>> Page.objects.search("about | documentation | django | home", raw=True)
[<Page: Page: Home page>, <Page: Page: About>, <Page: Page: Navigation>]
```

Django 1.10

```
>>> Entry.objects.filter(body_text__search='recipe')
[<Entry: Cheese on Toast recipes>, <Entry: Pizza
recipes>]

>>> Entry.objects.annotate(
...     search=SearchVector('blog__tagline',
'body_text'),
... ).filter(search='cheese')
[
    <Entry: Cheese on Toast recipes>,
    <Entry: Pizza Recipes>,
    <Entry: Dairy farming in Argentina>,
]
```

PostgreSQL FTS

Pros:

- Quick implementation
- No dependency

Cons:

- Need manually manage indexes
- depend on PostgreSQL
- no analytics data
- no DSL only `&` and `|` queries

ElasticSearch

The Lucene logo consists of the word "Lucene" in a stylized, green, wavy font. The letter "L" has a black feather-like flourish on its left side.



Who uses ElasticSearch?

GitHub



stackoverflow

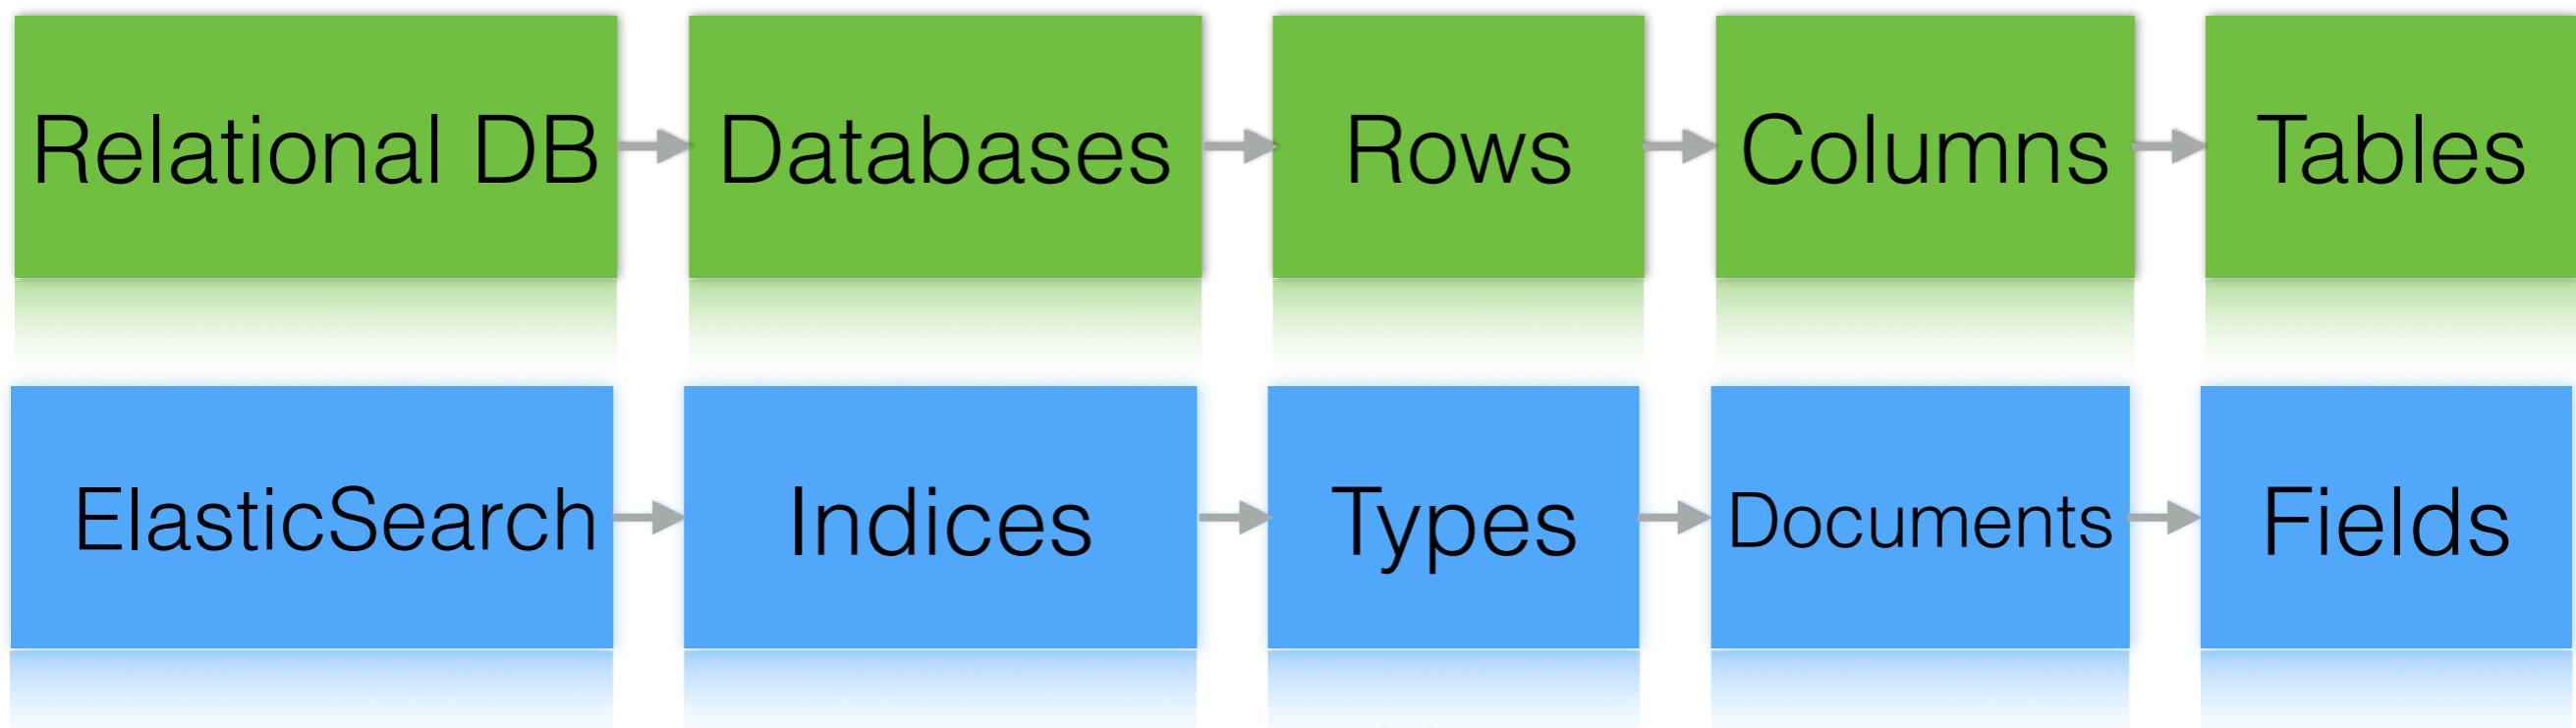


the guardian



WIKIPEDIA
The Free Encyclopedia

ElasticSearch: Quick Intro



ElasticSearch: Locks

- Pessimistic concurrency control
- Optimistic concurrency control

ElasticSearch and Python

- elasticsearch-py
- elasticsearch-py-async by Honza Kral
- elasticsearch-dsl-py by Honza Kral

ElasticSearch: FTS

```
$ curl -XGET 'http://localhost:9200/  
pyconua/talk/_search' -d '  
{  
    "query": {  
        "match": {  
            "user": "Andrii"  
        }  
    }  
}'
```

ES: Create Index

```
$ curl -XPUT 'http://localhost:9200/
twitter/' -d '{
  "settings" : {
    "index" : {
      "number_of_shards" : 3,
      "number_of_replicas" : 2
    }
  }
}'
```

ES: Add json to Index

```
$ curl -XPUT 'http://localhost:9200/  
pyconua/talk/1' -d '{  
  "user" : "andrii",  
  "description" : "Full text search"  
}'
```

ES: Stopwords

```
$ curl -XPUT 'http://localhost:9200/europython' -d
'{
  "settings": {
    "analysis": {
      "analyzer": {
        "my_english": {
          "type": "english",
          "stopwords_path": "stopwords/english.txt"
        }
      }
    }
  }
}'
```

ES: Highlight

```
$ curl -XGET 'http://localhost:9200/europython/talk/_search' -d '{  
  "query" : {...},  
  "highlight" : {  
    "pre_tags" : ["<tag1>"],  
    "post_tags" : ["</tag1>"],  
    "fields" : {  
      "_all" : {}  
    }  
  }'  
}'
```

ES: Relevance

```
$ curl -XGET 'http://localhost:9200/_search?explain -d
'
{
  "query" : { "match" : { "user" : "andrii" } }
}

"_explanation": {
  "description": "weight(tweet:honeymoon in 0)
                  [PerFieldSimilarity], result of:",
  "value":      0.076713204,
  "details": [...]
}
```



- written in C+
- uses MySQL as data source (or other database)

Sphinx search server

DB table ≈ Sphinx index

DB rows ≈ Sphinx documents

DB columns ≈ Sphinx fields and attributes

Sphinx simple query

```
SELECT * FROM test1
WHERE MATCH('europython');
```

Whoosh

- Pure-Python
- **Whoosh** was created by *Matt Chaput*.
- Pluggable scoring algorithm (including BM25F)
- more info at video from PyCon US 2013

Whoosh: Stop words

```
import os.path
import textwrap
```

```
names = os.listdir("stopwords")
for name in names:
    f = open("stopwords/" + name)
    wordls = [line.strip() for line in f]
    words = " ".join(wordls)
    print "%s": frozenset(u"""\n        """ % name
    print textwrap.fill(words, 72)
    print """\n        """.split())'
```

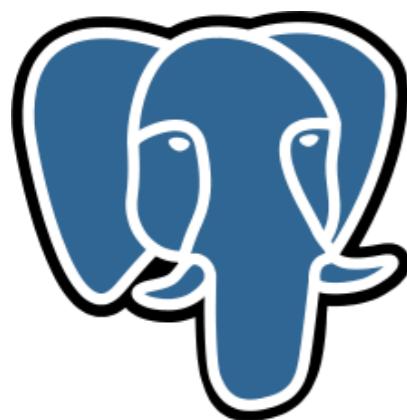
[http://anoncvs.postgresql.org/cvsweb.cgi/pgsql/src/backend/
snowball/stopwords/](http://anoncvs.postgresql.org/cvsweb.cgi/pgsql/src/backend/snowball/stopwords/)

Whoosh: Highlight

```
results = pycon.search(myquery)
for hit in results:
    print(hit["title"])
    # Assume "content" field is stored
    print(hit.highlights("content"))
```

Whoosh: Ranking search results

- Pluggable scoring algorithm
- including BM25F



Python clients

Python 3

Django support

elasticsearch-py
elasticsearch-dsl-py
elasticsearch-py-
async

YES

haystack +
elasticstack

psycopg2
aiopg
asyncpg

YES

djorm-ext-
pgfulltext
django.contrib.po
stgres

sphinxapi

NOT YET
(Open PR)

django-sphinx
django-sphinxql

Whoosh

YES

support using
haystack

Haystack



Haystack

Modular search for django

Haystack



Haystack: Pros and Cons

Pros:

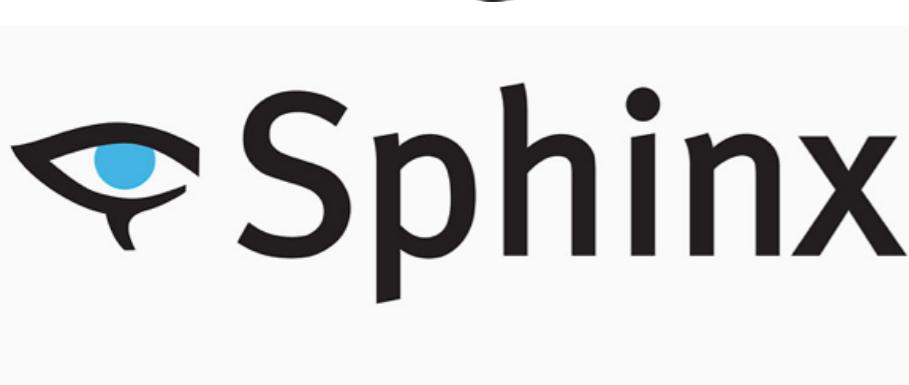
- easy to setup
- looks like Django ORM but for searches
- search engine independent
- support 4 engines (Elastic, Solr, Xapian, Whoosh)

Cons:

- poor SearchQuerySet API
- difficult to manage stop words
- loose performance, because extra layer
- Model - based

Indexes

Without indexes



Apache Lucene

No support

GIN / GIST

`to_tsvector()`

Disk / RT / Distributed

No support

index folder

No support

highlight
search
results

ranking /
relevance

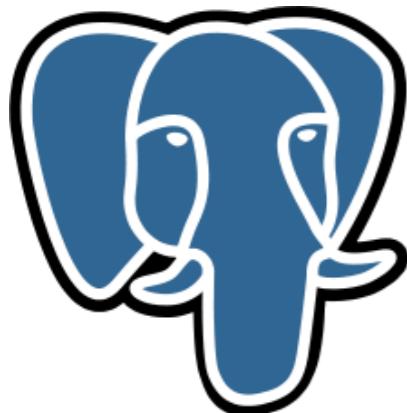
Configure
Stopwords



TF/IDF

YES

YES



cd_rank

YES

YES



max_lcs+BM25

YES

YES



Okapi BM25

YES

YES

Synonyms

Scale



elastic

YES

YES



YES

Partitioning



YES

Distributed searching



NO SUPPORT

NO

1 million music Artists

Evie Tamala

Jean-Pierre Martin

Deejay One

wecamewithbrokenteeth

The Blackbelt Band

Giant Tomo

Decoding Jesus

Elvin Jones & Jimmy Garrison Sextet

Infester

...

David Silverman

Aili Teigmo

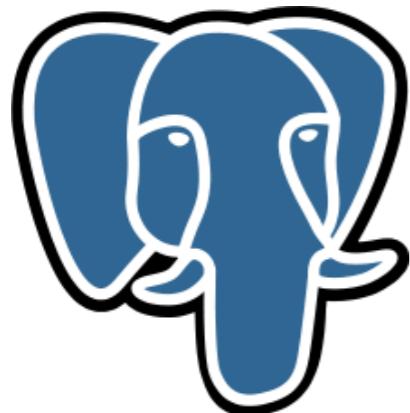
Performance

Database size



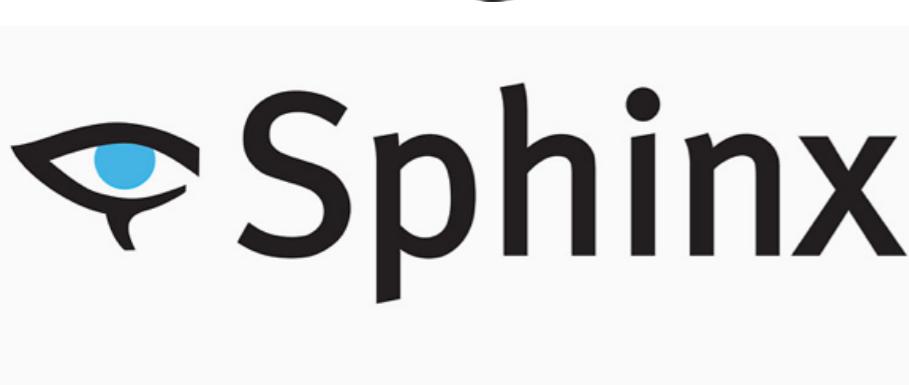
9 ms

~ 1 million records



4 ms

~ 1 million records



6 ms

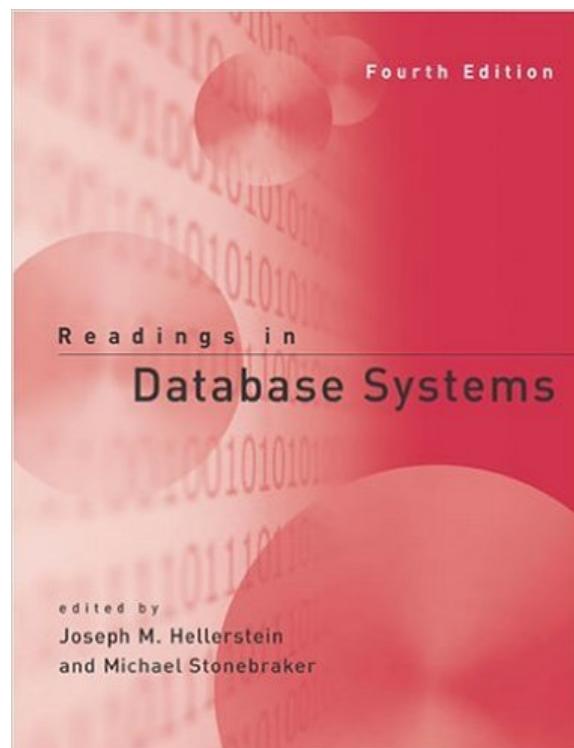
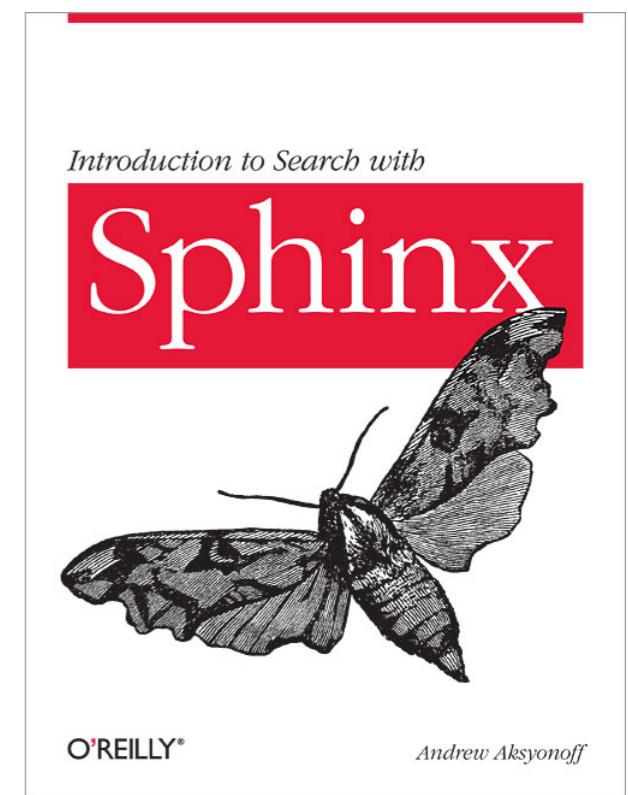
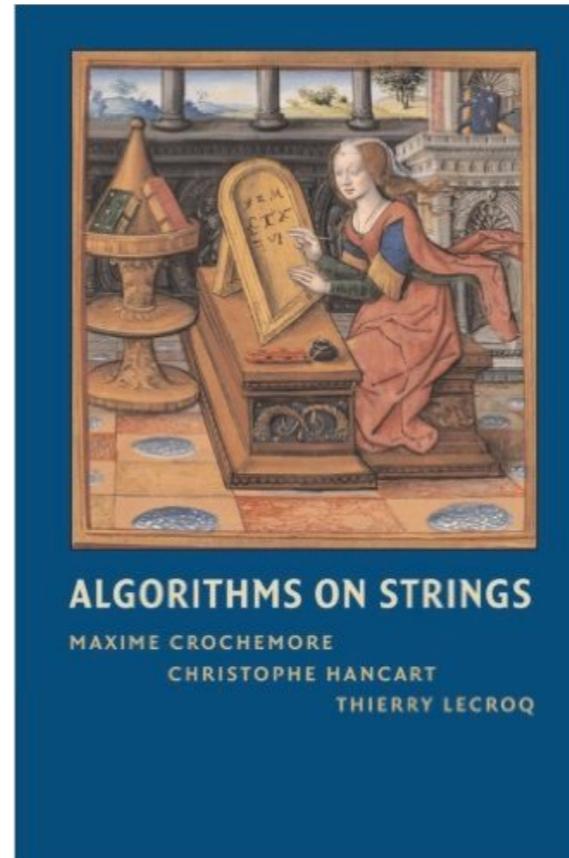
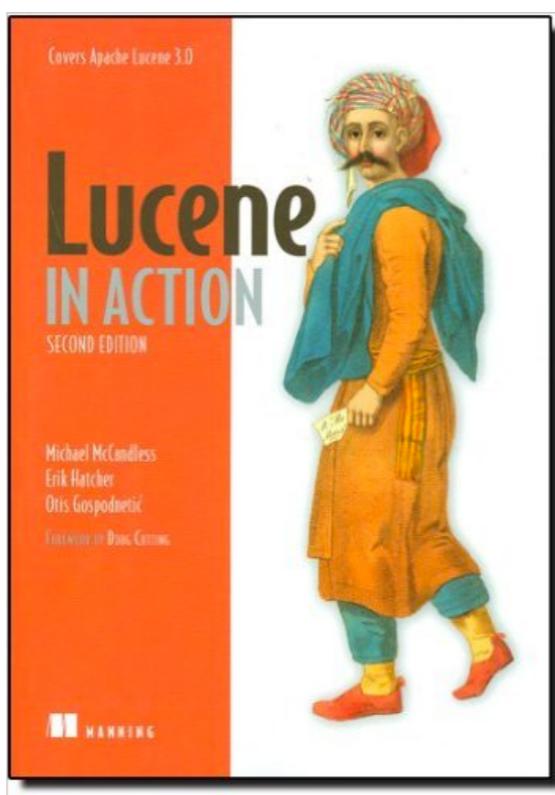
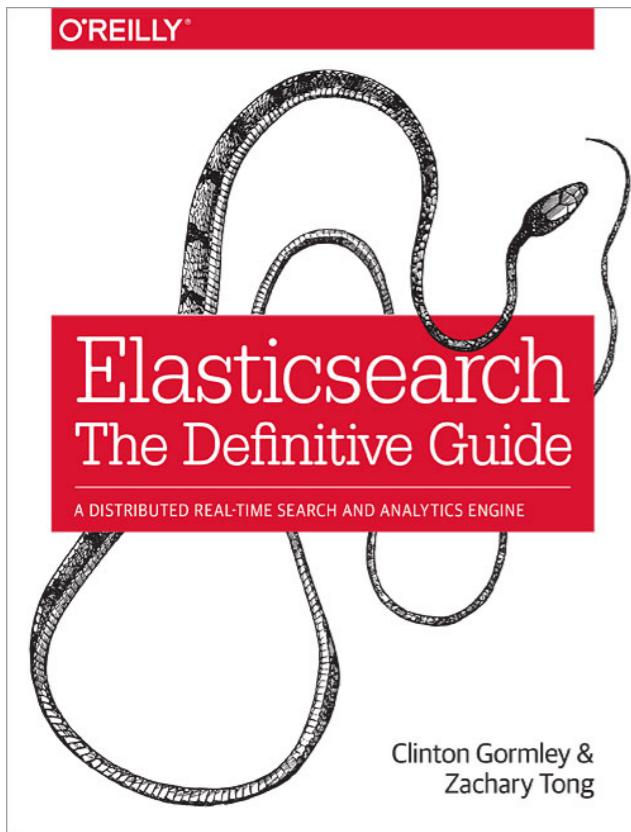
~ 1 million records



~2 s

~ 1 million records

Books



Indexing references:

<http://gist.cs.berkeley.edu/>

<http://www.sai.msu.su/~megera/postgres/gist/>

<http://www.sai.msu.su/~megera/wiki/Gin>

<https://www.postgresql.org/docs/9.5/static/gist.html>

<https://www.postgresql.org/docs/9.5/static/gin.html>

Ranking references:

<http://sphinxsearch.com/docs/current.html#weighting>

<https://www.postgresql.org/docs/9.5/static/textsearch-controls.html#TEXTSEARCH-RANKING>

<https://www.elastic.co/guide/en/elasticsearch/guide/current/scoring-theory.html>

https://en.wikipedia.org/wiki/Okapi_BM25

https://lucene.apache.org/core/3_6_0/scoring.html

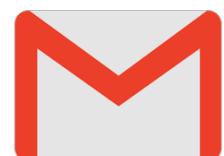
Slides

<https://asoldatenko.com/EuroPython16.pdf>

Thank You



@a_soldatenko



andrii.soldatenko@gmail.com



Hire the top 3% of freelance developers

<http://bit.ly/21lxQ01>



(n)->[:Questions]

